Pneumonia_Detection

November 11, 2024

1 Pneumonia Detection from Chest X-Ray Images

Dataset available on kaggle: https://www.kaggle.com/paultimothymooney/chest-xray-pneumonia

1.1 Install and load modules

[1]:
```
# !pip install timm torchsummary tqdm numpy matplotlib torch torchvision␣
↪kagglehub
```

[2]:
```
!pip install -q torchmetrics
```

```
                                        891.4/891.4 kB
18.4 MB/s eta 0:00:00
```

[3]:
```
import torch
import numpy as np
import matplotlib.pyplot as plt
import os
from tqdm.notebook import tqdm
import torchmetrics
```

1.2 Download the dataset

[4]:
```
import kagglehub

# Download latest version
path = kagglehub.dataset_download("paultimothymooney/chest-xray-pneumonia")

print("Path to dataset files:", path)
```

```
Downloading from
https://www.kaggle.com/api/v1/datasets/download/paultimothymooney/chest-xray-
pneumonia?dataset_version_number=2…

100%|     | 2.29G/2.29G [00:23<00:00, 104MB/s]

Extracting files…
```

```
Path to dataset files: /root/.cache/kagglehub/datasets/paultimothymooney/chest-
xray-pneumonia/versions/2
```

[5]:
```
!pwd
```

```
/content
```

[6]:
```
! cp -r /root/.cache/kagglehub/datasets/paultimothymooney/chest-xray-pneumonia/
 ↪versions/2 /content/
```

1.3 Define problem parameters

[7]:
```
class config:

    epochs = 1                                          # No. of epochs of training␣
 ↪the model
    lr_custom = 0.000005                                # Learning rate for custom␣
 ↪model
    lr_pretrained = 0.00001                             # Learning rate for␣
 ↪pretrained model
    batch_size = 128                                    # Batch Size For Dataset

    model_name = 'tf_efficientnet_b7.ap_in1k'       # Model name (We are going␣
 ↪to import model from timm)

    img_size = 224

    in_mean = [0.485, 0.456, 0.406]  # ImageNet dataset mean
    in_std = [0.229, 0.224, 0.225]   # ImageNet dataset std

    pn_mean = [0.4752]  # Pneumonia dataset mean
    pn_std = [0.2234]   # Pneumonia dataset std

    # Going to be use for loading dataset
    data_dir = '/content/2/chest_xray/chest_xray'    # Data Directory
    test_dir = 'test'                                   # Test folder name in data␣
 ↪directory
    train_dir = 'train'                                 # Train folder name in data␣
 ↪directory
    valid_dir = 'val'                                   # Valid folder name in data␣
 ↪directory
```

1.4 Choose available device

[8]:
```
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

print(f"device: {device}")

if torch.cuda.is_available():
    !nvidia-smi
```

```
device: cuda
Mon Nov 11 05:27:06 2024
+---------------------------------------------------------------------------------------
--------+
| NVIDIA-SMI 535.104.05             Driver Version: 535.104.05   CUDA Version:
12.2     |
|-----------------------------------------+----------------------+----------------------
--------+
| GPU  Name                 Persistence-M | Bus-Id        Disp.A | Volatile
Uncorr. ECC |
| Fan  Temp   Perf          Pwr:Usage/Cap |         Memory-Usage | GPU-Util
Compute M. |
|                                         |                      |
MIG M. |
|=========================================+======================+======================
========|
|   0  Tesla T4                       Off | 00000000:00:04.0 Off |
0 |
| N/A   51C    P8              10W /  70W |      3MiB / 15360MiB |      0%
Default |
|                                         |                      |
N/A |
+-----------------------------------------+----------------------+----------------------
--------+

+---------------------------------------------------------------------------------------
--------+
| Processes:
|
|  GPU   GI   CI        PID   Type   Process name                            GPU
Memory |
|        ID   ID
Usage      |
|=======================================================================================
========|
|  No running processes found
|
+---------------------------------------------------------------------------------------
--------+
```

1.5 Image transformation pipelines

[9]:
```
from torchvision import transforms as T, datasets
from torch.utils.data import Dataset
```

[10]:
```
class Gray2RGB:
    def __call__(self, image):
        return image.repeat(3, 1, 1) # Repeat the single channel across 3␣
 ↪channels to convert to RGB

train_transform_custom = T.Compose([
                            T.Resize(size=(config.img_size, config.img_size)),
                            T.RandomHorizontalFlip(p=0.5),
                            T.RandomRotation(degrees=(-20, 20)),
                            T.ToTensor(),
                            T.Normalize(config.pn_mean, config.pn_std),
                            Gray2RGB(),
                        ])

valid_transform_custom = T.Compose([
                            T.Resize(size=(config.img_size, config.img_size)),
                            T.ToTensor(),
                            T.Normalize(config.pn_mean, config.pn_std),
                            Gray2RGB(),
                        ])

test_transform_custom = T.Compose([
                            T.Resize(size=(config.img_size, config.img_size)),
                            T.ToTensor(),
                            T.Normalize(config.pn_mean, config.pn_std),
                            Gray2RGB(),
                        ])
```

[11]:
```
train_transform_pretrained = T.Compose([
                            T.Resize(size=(config.img_size, config.img_size)),
                            T.RandomHorizontalFlip(p=0.5),
                            T.RandomRotation(degrees=(-20, 20)),
                            T.ToTensor(),
                            Gray2RGB(),
                            T.Normalize(config.in_mean, config.in_std)
                        ])

valid_transform_pretrained = T.Compose([
                            T.Resize(size=(config.img_size, config.img_size)),
                            T.ToTensor(),
                            Gray2RGB(),
                            T.Normalize(config.in_mean, config.in_std)
```

```
                        ])

test_transform_pretrained = T.Compose([
                            T.Resize(size=(config.img_size, config.img_size)),
                            T.ToTensor(),
                            Gray2RGB(),
                            T.Normalize(config.in_mean, config.in_std)
                        ])
```

1.6 Pre-load datasets

[12]:
```
train_path = os.path.join(config.data_dir, config.train_dir)
valid_path = os.path.join(config.data_dir, config.valid_dir)
test_path = os.path.join(config.data_dir, config.test_dir)
```

[13]:
```
trainset_temp = datasets.ImageFolder(train_path, transform=None)
validset_temp = datasets.ImageFolder(valid_path, transform=None)
testset_temp = datasets.ImageFolder(test_path, transform=None)
```

[14]:
```
print(f"Train set size : {len(trainset_temp)}")
print(f"Valid set size : {len(validset_temp)}")
print(f"Test set size : {len(testset_temp)}")
```

```
Train set size : 5216
Valid set size : 16
Test set size : 624
```

[15]:
```
# Loading the datasets into the memory to boost performance

train_img_list = []
train_label_list = []
for img, label in trainset_temp:
    img = img.convert("L")
    train_img_list.append(img)
    train_label_list.append(label)

valid_img_list = []
valid_label_list = []
for img, label in validset_temp:
    img = img.convert("L")
    valid_img_list.append(img)
    valid_label_list.append(label)

test_img_list = []
test_label_list = []
for img, label in testset_temp:
    img = img.convert("L")
```

```
    test_img_list.append(img)
    test_label_list.append(label)
```

1.7 Create custom dataset

[16]:
```
class CustomDataset(Dataset):

    def __init__(self, img_list, label_list, transform=None):
        super().__init__()
        self.img_list = img_list
        self.label_list = label_list
        self.transform = transform

    def __len__(self):
        return len(self.img_list)

    def __getitem__(self, index):
        image = self.img_list[index]
        label = self.label_list[index]

        if self.transform:
            image = self.transform(image)

        return image, label
```

[17]:
```
trainset_custom = CustomDataset(train_img_list, train_label_list,␣
 ↪train_transform_custom)
trainset_pretrained = CustomDataset(train_img_list, train_label_list,␣
 ↪train_transform_pretrained)

validset_custom = CustomDataset(valid_img_list, valid_label_list,␣
 ↪valid_transform_custom)
validset_pretrained = CustomDataset(valid_img_list, valid_label_list,␣
 ↪valid_transform_pretrained)

testset_custom = CustomDataset(test_img_list, test_label_list,␣
 ↪test_transform_custom)
testset_pretrained = CustomDataset(test_img_list, test_label_list,␣
 ↪test_transform_pretrained)
```

1.8 Count samples in each class

[18]:
```
from collections import Counter

# Function for counting the number of samples in each class
def count_classes(dataset):
```

```
    label_counts = Counter()
    for _, label in dataset:
        label_counts[label] += 1
    return label_counts
```

[19]:
```
label_counts_train = count_classes(trainset_custom)

print(label_counts_train)  # Counter({1: 3875, 0: 1341})
```

```
Counter({1: 3875, 0: 1341})
```

[20]:
```
label_counts_valid = count_classes(validset_custom)

print(label_counts_valid)
```

```
Counter({0: 8, 1: 8})
```

[21]:
```
label_counts_test = count_classes(testset_custom)

print(label_counts_test)
```

```
Counter({1: 390, 0: 234})
```

[22]:
```
sizes = [label_counts_train[0], label_counts_train[1]]
labels = ['Normal', 'Pneumonia']

plt.pie(sizes, labels=labels, autopct='%1.1f%%');
```

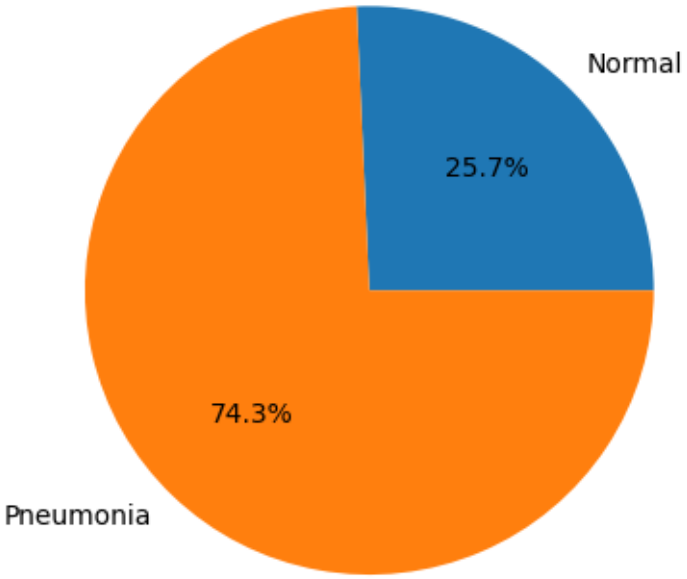

1.9 Calculate the weight for each class

[23]:
```
# calculating weight for each class
from sklearn.utils.class_weight import compute_class_weight

# Get class labels from dataset
y_labels = [label for _, label in trainset_custom]
class_weights = compute_class_weight(class_weight='balanced', classes=np.
 ↪unique(y_labels), y=y_labels)

# Convert to torch tensor
class_weights = torch.tensor(class_weights, dtype=torch.float)

print(f'class_weights: {class_weights}')
```

```
class_weights: tensor([1.9448, 0.6730])
```

[24]:
```
def show_image(image, label='-', mean=None, std=None):

    image = image.permute(1, 2, 0)

    if mean and std:
        mean = torch.FloatTensor(mean)
```

```
        std = torch.FloatTensor(std)
        image = image*std + mean
        image = torch.clip(image, 0, 1)
        image = image.numpy()

    plt.imshow(image, cmap="gray")
    plt.title(label)
```

[25]:
```
sample_img, sample_label = trainset_custom[2]

show_image(sample_img, labels[sample_label], config.pn_mean, config.pn_std)
```

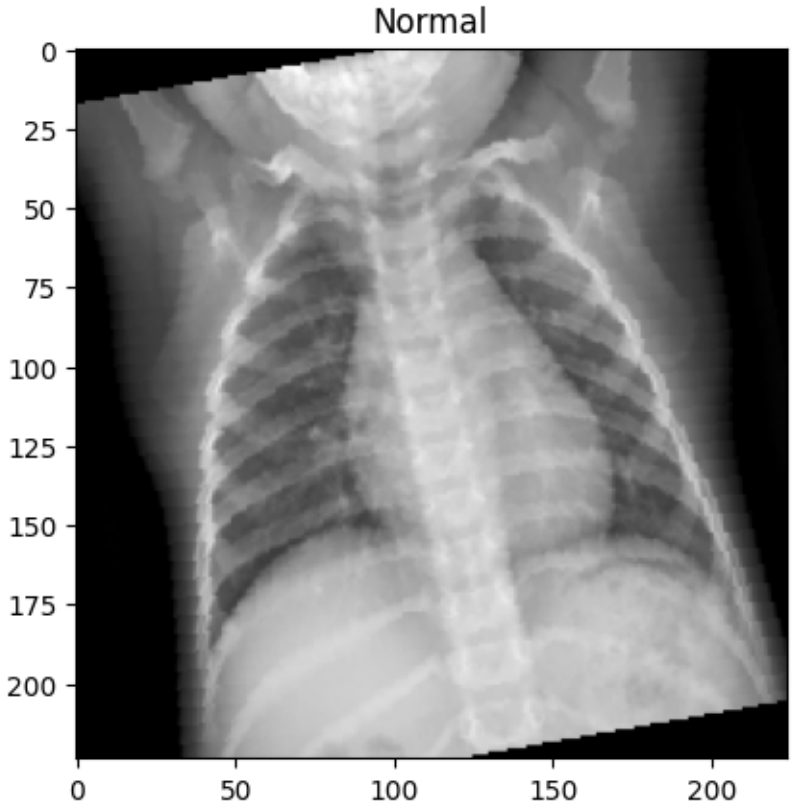

1.10 Load dataset into batches

[26]:
```
from torch.utils.data import DataLoader
from torchvision.utils import make_grid
```

[27]:
```
train_loader_custom = DataLoader(trainset_custom, batch_size=config.batch_size,␣
↪shuffle=True)
```

```
valid_loader_custom = DataLoader(validset_custom, batch_size=config.batch_size,␣
 ↪shuffle=False)
test_loader_custom = DataLoader(testset_custom, batch_size=config.batch_size,␣
 ↪shuffle=False)
```

[28]:
```
train_loader_pretrained = DataLoader(trainset_pretrained, batch_size=config.
 ↪batch_size, shuffle=True)
valid_loader_pretrained = DataLoader(validset_pretrained, batch_size=config.
 ↪batch_size, shuffle=False)
test_loader_pretrained = DataLoader(testset_pretrained, batch_size=config.
 ↪batch_size, shuffle=False)
```

[29]:
```
print(f"No. of batches in train loader : {len(train_loader_custom)}")
print(f"No. of batches in valid loader : {len(valid_loader_custom)}")
print(f"No. of batches in test loader : {len(test_loader_custom)}")
```

```
No. of batches in train loader : 41
No. of batches in valid loader : 1
No. of batches in test loader : 5
```

[30]:
```
train_iter = iter(train_loader_custom)
sample_batch, sample_batch_label = next(train_iter)

print(f'sample_batch.shape: {sample_batch.shape}')
```

```
sample_batch.shape: torch.Size([128, 3, 224, 224])
```

1.11 Show samples of dataset

[31]:
```
fig, ax = plt.subplots(2, 8, figsize=(12,3))

sample_batch, sample_batch_label = next(iter(train_loader_custom))
for n in range(2):
    for m in range(8):
        image = sample_batch[n*8+m].cpu().permute(1, 2, 0)
        mean = torch.FloatTensor(config.pn_mean)
        std = torch.FloatTensor(config.pn_std)
        image = image*std + mean
        image = torch.clip(image, 0, 1)
        image = image.numpy()

        ax[n,m].imshow(image, cmap="gray")
        ax[n,m].grid(False)
        ax[n,m].axis(False)
        ax[n,m].tick_params(labelbottom=False, labelleft=False)
        ax[n,m].set_title(f'{labels[sample_batch_label[n*8+m]]}')
```

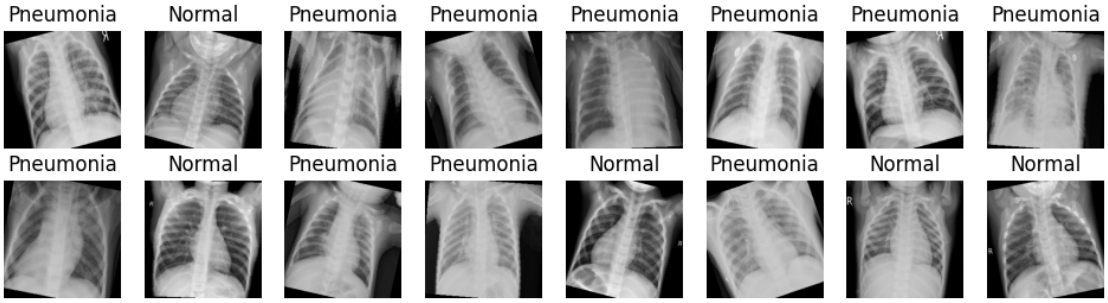

1.12 Design custom CNN model

[32]:
```
import torch.nn as nn
import torch.nn.functional as F
```

[33]:
```
# Calculating the size of the final layer before feeding into the linear layer
X = torch.zeros(1, 3, config.img_size, config.img_size)

conv1 = nn.Sequential(
                nn.Conv2d(3, 8, 3, stride=1, padding=1),
                nn.BatchNorm2d(8),
                nn.ReLU(inplace=True),
                nn.MaxPool2d(2,2))

conv2 = nn.Sequential(
                nn.Conv2d(8, 16, 3, stride=1, padding=1),
                nn.BatchNorm2d(16),
                nn.ReLU(inplace=True),
                nn.MaxPool2d(2,2))

conv3 = nn.Sequential(
                nn.Conv2d(16, 32, 3, stride=1, padding=1),
                nn.BatchNorm2d(32),
                nn.ReLU(inplace=True),
                nn.MaxPool2d(2,2))

conv4 = nn.Sequential(
                nn.Conv2d(32, 64, 3, stride=1, padding=1),
                nn.BatchNorm2d(64),
                nn.ReLU(inplace=True),
                nn.MaxPool2d(2,2))

conv5 = nn.Sequential(
                nn.Conv2d(64, 128, 3, stride=1, padding=1),
                nn.BatchNorm2d(128),
```

```
                nn.ReLU(inplace=True),
                nn.MaxPool2d(2,2))

X = conv1(X)
X = conv2(X)
X = conv3(X)
X = conv4(X)
X = conv5(X)

X.shape # [1, 128, 7, 7]
```

[33]:
```
torch.Size([1, 128, 7, 7])
```

[34]:
```
class ConvolutionalNetwork(nn.Module):
    def __init__(self):
        super().__init__()

        self.conv1 = nn.Sequential(
                        nn.Conv2d(3, 8, 3, stride=1, padding=1),
                        nn.ReLU(inplace=True),
                        nn.BatchNorm2d(8),
                        nn.MaxPool2d(2,2))

        self.conv2 = nn.Sequential(
                        nn.Conv2d(8, 16, 3, stride=1, padding=1),
                        nn.ReLU(inplace=True),
                        nn.BatchNorm2d(16),
                        nn.MaxPool2d(2,2))

        self.conv3 = nn.Sequential(
                        nn.Conv2d(16, 32, 3, stride=1, padding=1),
                        nn.ReLU(inplace=True),
                        nn.BatchNorm2d(32),
                        nn.MaxPool2d(2,2))

        self.conv4 = nn.Sequential(
                        nn.Conv2d(32, 64, 3, stride=1, padding=1),
                        nn.ReLU(inplace=True),
                        nn.BatchNorm2d(64),
                        nn.MaxPool2d(2,2))

        self.conv5 = nn.Sequential(
                        nn.Conv2d(64, 128, 3, stride=1, padding=1),
                        nn.ReLU(inplace=True),
                        nn.BatchNorm2d(128),
                        nn.MaxPool2d(2,2))
```

```
        self.fc = nn.Sequential(
                    nn.Linear(128*7*7, 512),
                    nn.ReLU(inplace=True),
                    nn.BatchNorm1d(512),
                    nn.Dropout(0.5),
                    nn.Linear(512, 2))

    def forward(self, x):
        x = self.conv1(x)
        x = self.conv2(x)
        x = self.conv3(x)
        x = self.conv4(x)
        x = self.conv5(x)
        x = x.view(x.shape[0], -1)
        x = self.fc(x)
        return x
```

[35]:
```
torch.manual_seed(42)
torch.cuda.manual_seed(42)
cnn_model = ConvolutionalNetwork()

cnn_model.to(device)
```

[35]:
```
ConvolutionalNetwork(
  (conv1): Sequential(
    (0): Conv2d(3, 8, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace=True)
    (2): BatchNorm2d(8, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
  )
  (conv2): Sequential(
    (0): Conv2d(8, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace=True)
    (2): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
  )
  (conv3): Sequential(
    (0): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace=True)
    (2): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
```

```
  )
  (conv4): Sequential(
    (0): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace=True)
    (2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
  )
  (conv5): Sequential(
    (0): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace=True)
    (2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
  )
  (fc): Sequential(
    (0): Linear(in_features=6272, out_features=512, bias=True)
    (1): ReLU(inplace=True)
    (2): BatchNorm1d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (3): Dropout(p=0.5, inplace=False)
    (4): Linear(in_features=512, out_features=2, bias=True)
  )
)
```

[36]:
```
from torchsummary import summary

summary(cnn_model, input_size=(3, config.img_size, config.img_size))
```

```
----------------------------------------------------------------
        Layer (type)               Output Shape         Param #
================================================================
            Conv2d-1          [-1, 8, 224, 224]             224
              ReLU-2          [-1, 8, 224, 224]               0
       BatchNorm2d-3          [-1, 8, 224, 224]              16
         MaxPool2d-4          [-1, 8, 112, 112]               0
            Conv2d-5         [-1, 16, 112, 112]           1,168
              ReLU-6         [-1, 16, 112, 112]               0
       BatchNorm2d-7         [-1, 16, 112, 112]              32
         MaxPool2d-8           [-1, 16, 56, 56]               0
            Conv2d-9           [-1, 32, 56, 56]           4,640
             ReLU-10           [-1, 32, 56, 56]               0
      BatchNorm2d-11           [-1, 32, 56, 56]              64
        MaxPool2d-12           [-1, 32, 28, 28]               0
           Conv2d-13           [-1, 64, 28, 28]          18,496
```

```
              ReLU-14           [-1, 64, 28, 28]               0
       BatchNorm2d-15           [-1, 64, 28, 28]             128
         MaxPool2d-16           [-1, 64, 14, 14]               0
            Conv2d-17          [-1, 128, 14, 14]          73,856
              ReLU-18          [-1, 128, 14, 14]               0
       BatchNorm2d-19          [-1, 128, 14, 14]             256
         MaxPool2d-20            [-1, 128, 7, 7]               0
            Linear-21                  [-1, 512]       3,211,776
              ReLU-22                  [-1, 512]               0
       BatchNorm1d-23                  [-1, 512]           1,024
           Dropout-24                  [-1, 512]               0
            Linear-25                    [-1, 2]           1,026
================================================================
Total params: 3,312,706
Trainable params: 3,312,706
Non-trainable params: 0
----------------------------------------------------------------
Input size (MB): 0.57
Forward/backward pass size (MB): 19.30
Params size (MB): 12.64
Estimated Total Size (MB): 32.51
----------------------------------------------------------------
```

1.13 Define pretrained models for fine-tuning

[37]:
```python
from torch import nn
import torch.nn.functional as F
import timm

# Models that we will be using:
# tf_efficientnet_b4.ap_in1k
# tf_efficientnet_b7.ap_in1k
# resnet50.tv_in1k
# resnet101.tv_in1k

torch.manual_seed(42)
torch.cuda.manual_seed(42)

config.model_name = 'tf_efficientnet_b7.ap_in1k'

model = timm.create_model(config.model_name, pretrained=True)
model
```

```
/usr/local/lib/python3.10/dist-packages/huggingface_hub/utils/_token.py:89:
UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab
(https://huggingface.co/settings/tokens), set it as secret in your Google Colab
```

```
and restart your session.
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access
public models or datasets.
  warnings.warn(

model.safetensors:   0%|          | 0.00/267M [00:00<?, ?B/s]
```

[37]:
```
EfficientNet(
  (conv_stem): Conv2dSame(3, 64, kernel_size=(3, 3), stride=(2, 2), bias=False)
  (bn1): BatchNormAct2d(
    64, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
    (drop): Identity()
    (act): SiLU(inplace=True)
  )
  (blocks): Sequential(
    (0): Sequential(
      (0): DepthwiseSeparableConv(
        (conv_dw): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=64, bias=False)
        (bn1): BatchNormAct2d(
          64, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(64, 16, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(16, 64, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pw): Conv2d(64, 32, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn2): BatchNormAct2d(
          32, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (1): DepthwiseSeparableConv(
        (conv_dw): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=32, bias=False)
        (bn1): BatchNormAct2d(
          32, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
```

```
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(32, 8, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(8, 32, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pw): Conv2d(32, 32, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn2): BatchNormAct2d(
          32, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (2): DepthwiseSeparableConv(
        (conv_dw): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=32, bias=False)
        (bn1): BatchNormAct2d(
          32, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(32, 8, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(8, 32, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pw): Conv2d(32, 32, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn2): BatchNormAct2d(
          32, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (3): DepthwiseSeparableConv(
        (conv_dw): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=32, bias=False)
        (bn1): BatchNormAct2d(
          32, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
```

```
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(32, 8, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(8, 32, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pw): Conv2d(32, 32, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn2): BatchNormAct2d(
          32, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
    )
    (1): Sequential(
      (0): InvertedResidual(
        (conv_pw): Conv2d(32, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2dSame(192, 192, kernel_size=(3, 3), stride=(2, 2),
groups=192, bias=False)
        (bn2): BatchNormAct2d(
          192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(192, 8, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(8, 192, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(192, 48, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          48, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
```

```
      (1): InvertedResidual(
        (conv_pw): Conv2d(48, 288, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(288, 288, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=288, bias=False)
        (bn2): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(288, 12, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(12, 288, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(288, 48, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          48, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (2): InvertedResidual(
        (conv_pw): Conv2d(48, 288, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(288, 288, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=288, bias=False)
        (bn2): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
```

```
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(288, 12, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(12, 288, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(288, 48, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          48, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (3): InvertedResidual(
        (conv_pw): Conv2d(48, 288, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(288, 288, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=288, bias=False)
        (bn2): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(288, 12, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(12, 288, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(288, 48, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          48, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (4): InvertedResidual(
```

```
        (conv_pw): Conv2d(48, 288, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(288, 288, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=288, bias=False)
        (bn2): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(288, 12, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(12, 288, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(288, 48, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          48, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (5): InvertedResidual(
        (conv_pw): Conv2d(48, 288, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(288, 288, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=288, bias=False)
        (bn2): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
```

```
          (conv_reduce): Conv2d(288, 12, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(12, 288, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(288, 48, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          48, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (6): InvertedResidual(
        (conv_pw): Conv2d(48, 288, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(288, 288, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=288, bias=False)
        (bn2): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(288, 12, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(12, 288, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(288, 48, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          48, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
    )
    (2): Sequential(
```

```
      (0): InvertedResidual(
        (conv_pw): Conv2d(48, 288, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2dSame(288, 288, kernel_size=(5, 5), stride=(2, 2),
groups=288, bias=False)
        (bn2): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(288, 12, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(12, 288, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(288, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          80, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (1): InvertedResidual(
        (conv_pw): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(480, 480, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=480, bias=False)
        (bn2): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
```

```
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          80, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (2): InvertedResidual(
        (conv_pw): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(480, 480, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=480, bias=False)
        (bn2): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          80, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (3): InvertedResidual(
```

```
        (conv_pw): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(480, 480, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=480, bias=False)
        (bn2): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          80, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (4): InvertedResidual(
        (conv_pw): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(480, 480, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=480, bias=False)
        (bn2): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
```

```
          (conv_reduce): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          80, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (5): InvertedResidual(
        (conv_pw): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(480, 480, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=480, bias=False)
        (bn2): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          80, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (6): InvertedResidual(
        (conv_pw): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
```

```
bias=False)
        (bn1): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(480, 480, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=480, bias=False)
        (bn2): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          80, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
    )
    (3): Sequential(
      (0): InvertedResidual(
        (conv_pw): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2dSame(480, 480, kernel_size=(3, 3), stride=(2, 2),
groups=480, bias=False)
        (bn2): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
```

```
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(480, 160, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (1): InvertedResidual(
        (conv_pw): Conv2d(160, 960, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(960, 960, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=960, bias=False)
        (bn2): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(960, 40, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(40, 960, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(960, 160, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (2): InvertedResidual(
```

```
        (conv_pw): Conv2d(160, 960, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(960, 960, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=960, bias=False)
        (bn2): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(960, 40, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(40, 960, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(960, 160, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (3): InvertedResidual(
        (conv_pw): Conv2d(160, 960, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(960, 960, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=960, bias=False)
        (bn2): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
```

```
          (conv_reduce): Conv2d(960, 40, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(40, 960, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(960, 160, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (4): InvertedResidual(
        (conv_pw): Conv2d(160, 960, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(960, 960, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=960, bias=False)
        (bn2): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(960, 40, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(40, 960, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(960, 160, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (5): InvertedResidual(
        (conv_pw): Conv2d(160, 960, kernel_size=(1, 1), stride=(1, 1),
```

```
bias=False)
        (bn1): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(960, 960, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=960, bias=False)
        (bn2): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(960, 40, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(40, 960, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(960, 160, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (6): InvertedResidual(
        (conv_pw): Conv2d(160, 960, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(960, 960, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=960, bias=False)
        (bn2): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(960, 40, kernel_size=(1, 1), stride=(1, 1))
```

```
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(40, 960, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(960, 160, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (7): InvertedResidual(
        (conv_pw): Conv2d(160, 960, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(960, 960, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=960, bias=False)
        (bn2): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(960, 40, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(40, 960, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(960, 160, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (8): InvertedResidual(
        (conv_pw): Conv2d(160, 960, kernel_size=(1, 1), stride=(1, 1),
bias=False)
```

```
        (bn1): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(960, 960, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=960, bias=False)
        (bn2): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(960, 40, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(40, 960, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(960, 160, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (9): InvertedResidual(
        (conv_pw): Conv2d(160, 960, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(960, 960, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=960, bias=False)
        (bn2): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(960, 40, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
```

```
          (conv_expand): Conv2d(40, 960, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(960, 160, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
    )
    (4): Sequential(
      (0): InvertedResidual(
        (conv_pw): Conv2d(160, 960, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(960, 960, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=960, bias=False)
        (bn2): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(960, 40, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(40, 960, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(960, 224, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          224, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (1): InvertedResidual(
        (conv_pw): Conv2d(224, 1344, kernel_size=(1, 1), stride=(1, 1),
```

```
bias=False)
        (bn1): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(1344, 1344, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1344, bias=False)
        (bn2): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(1344, 56, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(56, 1344, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(1344, 224, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          224, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (2): InvertedResidual(
        (conv_pw): Conv2d(224, 1344, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(1344, 1344, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1344, bias=False)
        (bn2): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(1344, 56, kernel_size=(1, 1), stride=(1, 1))
```

```
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(56, 1344, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(1344, 224, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          224, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (3): InvertedResidual(
        (conv_pw): Conv2d(224, 1344, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(1344, 1344, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1344, bias=False)
        (bn2): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(1344, 56, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(56, 1344, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(1344, 224, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          224, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (4): InvertedResidual(
        (conv_pw): Conv2d(224, 1344, kernel_size=(1, 1), stride=(1, 1),
bias=False)
```

```
        (bn1): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(1344, 1344, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1344, bias=False)
        (bn2): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(1344, 56, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(56, 1344, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(1344, 224, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          224, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (5): InvertedResidual(
        (conv_pw): Conv2d(224, 1344, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(1344, 1344, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1344, bias=False)
        (bn2): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(1344, 56, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
```

```
          (conv_expand): Conv2d(56, 1344, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(1344, 224, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          224, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (6): InvertedResidual(
        (conv_pw): Conv2d(224, 1344, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(1344, 1344, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1344, bias=False)
        (bn2): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(1344, 56, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(56, 1344, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(1344, 224, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          224, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (7): InvertedResidual(
        (conv_pw): Conv2d(224, 1344, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
```

```
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(1344, 1344, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1344, bias=False)
        (bn2): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(1344, 56, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(56, 1344, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(1344, 224, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          224, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (8): InvertedResidual(
        (conv_pw): Conv2d(224, 1344, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(1344, 1344, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1344, bias=False)
        (bn2): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(1344, 56, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(56, 1344, kernel_size=(1, 1), stride=(1, 1))
```

```
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(1344, 224, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          224, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (9): InvertedResidual(
        (conv_pw): Conv2d(224, 1344, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(1344, 1344, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1344, bias=False)
        (bn2): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(1344, 56, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(56, 1344, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(1344, 224, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          224, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
    )
    (5): Sequential(
      (0): InvertedResidual(
        (conv_pw): Conv2d(224, 1344, kernel_size=(1, 1), stride=(1, 1),
bias=False)
```

```
        (bn1): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2dSame(1344, 1344, kernel_size=(5, 5), stride=(2, 2),
groups=1344, bias=False)
        (bn2): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(1344, 56, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(56, 1344, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(1344, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (1): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
```

```
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(2304, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (2): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(2304, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (3): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
```

```
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(2304, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (4): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
```

```
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(2304, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (5): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(2304, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (6): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
```

```
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(2304, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (7): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
```

```
        )
        (conv_pwl): Conv2d(2304, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (8): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(2304, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (9): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
```

```
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(2304, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (10): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
```

```
        (conv_pwl): Conv2d(2304, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (11): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(2304, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (12): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
```

```
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(2304, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
    )
    (6): Sequential(
      (0): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
```

```
        )
        (conv_pwl): Conv2d(2304, 640, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          640, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (1): InvertedResidual(
        (conv_pw): Conv2d(640, 3840, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          3840, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(3840, 3840, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=3840, bias=False)
        (bn2): BatchNormAct2d(
          3840, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(3840, 160, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(160, 3840, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(3840, 640, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          640, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (2): InvertedResidual(
        (conv_pw): Conv2d(640, 3840, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          3840, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
```

```
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(3840, 3840, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=3840, bias=False)
        (bn2): BatchNormAct2d(
          3840, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(3840, 160, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(160, 3840, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(3840, 640, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          640, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (3): InvertedResidual(
        (conv_pw): Conv2d(640, 3840, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          3840, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(3840, 3840, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=3840, bias=False)
        (bn2): BatchNormAct2d(
          3840, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(3840, 160, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(160, 3840, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
```

```
        (conv_pwl): Conv2d(3840, 640, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          640, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
    )
  )
  (conv_head): Conv2d(640, 2560, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn2): BatchNormAct2d(
    2560, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
    (drop): Identity()
    (act): SiLU(inplace=True)
  )
  (global_pool): SelectAdaptivePool2d(pool_type=avg,
flatten=Flatten(start_dim=1, end_dim=-1))
  (classifier): Linear(in_features=2560, out_features=1000, bias=True)
)
```

[38]:
```
for param in model.parameters():
    param.requires_grad = False

# Number of in- and out- features for each model:
# 1792 -> 768  -> 256 -> 2   EfficientNet_b4
# 2560 -> 1024 -> 256 -> 2   EfficientNet_b7
# 2048 -> 1024 -> 256 -> 2   ResNet50
# 2048 -> 1024 -> 256 -> 2   ResNet101

# model.classifier for efficientnet family
# model.fc for ResNet family
model.classifier = nn.Sequential(
    nn.Linear(in_features=2560, out_features=1024),
    nn.ReLU(),
    nn.Dropout(p=0.4),
    nn.Linear(in_features=1024, out_features=256),
    nn.ReLU(),
    nn.Linear(in_features=256, out_features=2),
)

model.to(device)
```

[38]:
```
EfficientNet(
  (conv_stem): Conv2dSame(3, 64, kernel_size=(3, 3), stride=(2, 2), bias=False)
  (bn1): BatchNormAct2d(
```

```
    64, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
    (drop): Identity()
    (act): SiLU(inplace=True)
  )
  (blocks): Sequential(
    (0): Sequential(
      (0): DepthwiseSeparableConv(
        (conv_dw): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=64, bias=False)
        (bn1): BatchNormAct2d(
          64, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(64, 16, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(16, 64, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pw): Conv2d(64, 32, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn2): BatchNormAct2d(
          32, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (1): DepthwiseSeparableConv(
        (conv_dw): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=32, bias=False)
        (bn1): BatchNormAct2d(
          32, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(32, 8, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(8, 32, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pw): Conv2d(32, 32, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn2): BatchNormAct2d(
          32, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
```

```
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (2): DepthwiseSeparableConv(
        (conv_dw): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=32, bias=False)
        (bn1): BatchNormAct2d(
          32, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(32, 8, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(8, 32, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pw): Conv2d(32, 32, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn2): BatchNormAct2d(
          32, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (3): DepthwiseSeparableConv(
        (conv_dw): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=32, bias=False)
        (bn1): BatchNormAct2d(
          32, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(32, 8, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(8, 32, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pw): Conv2d(32, 32, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn2): BatchNormAct2d(
          32, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
```

```
          (act): Identity()
        )
        (drop_path): Identity()
      )
    )
    (1): Sequential(
      (0): InvertedResidual(
        (conv_pw): Conv2d(32, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2dSame(192, 192, kernel_size=(3, 3), stride=(2, 2),
groups=192, bias=False)
        (bn2): BatchNormAct2d(
          192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(192, 8, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(8, 192, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(192, 48, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          48, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (1): InvertedResidual(
        (conv_pw): Conv2d(48, 288, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(288, 288, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=288, bias=False)
```

```
        (bn2): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(288, 12, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(12, 288, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(288, 48, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          48, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (2): InvertedResidual(
        (conv_pw): Conv2d(48, 288, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(288, 288, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=288, bias=False)
        (bn2): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(288, 12, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(12, 288, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(288, 48, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          48, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
```

```
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (3): InvertedResidual(
        (conv_pw): Conv2d(48, 288, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(288, 288, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=288, bias=False)
        (bn2): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(288, 12, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(12, 288, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(288, 48, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          48, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (4): InvertedResidual(
        (conv_pw): Conv2d(48, 288, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(288, 288, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=288, bias=False)
        (bn2): BatchNormAct2d(
```

```
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(288, 12, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(12, 288, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(288, 48, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          48, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (5): InvertedResidual(
        (conv_pw): Conv2d(48, 288, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(288, 288, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=288, bias=False)
        (bn2): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(288, 12, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(12, 288, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(288, 48, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          48, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
```

```
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (6): InvertedResidual(
        (conv_pw): Conv2d(48, 288, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(288, 288, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=288, bias=False)
        (bn2): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(288, 12, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(12, 288, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(288, 48, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          48, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
    )
    (2): Sequential(
      (0): InvertedResidual(
        (conv_pw): Conv2d(48, 288, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2dSame(288, 288, kernel_size=(5, 5), stride=(2, 2),
groups=288, bias=False)
```

```
        (bn2): BatchNormAct2d(
          288, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(288, 12, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(12, 288, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(288, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          80, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (1): InvertedResidual(
        (conv_pw): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(480, 480, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=480, bias=False)
        (bn2): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          80, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
```

```
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (2): InvertedResidual(
        (conv_pw): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(480, 480, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=480, bias=False)
        (bn2): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          80, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (3): InvertedResidual(
        (conv_pw): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(480, 480, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=480, bias=False)
        (bn2): BatchNormAct2d(
```

```
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          80, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (4): InvertedResidual(
        (conv_pw): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(480, 480, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=480, bias=False)
        (bn2): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          80, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
```

```
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (5): InvertedResidual(
        (conv_pw): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(480, 480, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=480, bias=False)
        (bn2): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          80, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (6): InvertedResidual(
        (conv_pw): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(480, 480, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=480, bias=False)
        (bn2): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
```

```
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          80, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
    )
    (3): Sequential(
      (0): InvertedResidual(
        (conv_pw): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2dSame(480, 480, kernel_size=(3, 3), stride=(2, 2),
groups=480, bias=False)
        (bn2): BatchNormAct2d(
          480, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(480, 160, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
```

```
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (1): InvertedResidual(
        (conv_pw): Conv2d(160, 960, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(960, 960, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=960, bias=False)
        (bn2): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(960, 40, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(40, 960, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(960, 160, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (2): InvertedResidual(
        (conv_pw): Conv2d(160, 960, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(960, 960, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=960, bias=False)
        (bn2): BatchNormAct2d(
```

```
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(960, 40, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(40, 960, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(960, 160, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (3): InvertedResidual(
        (conv_pw): Conv2d(160, 960, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(960, 960, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=960, bias=False)
        (bn2): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(960, 40, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(40, 960, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(960, 160, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
```

```
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (4): InvertedResidual(
        (conv_pw): Conv2d(160, 960, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(960, 960, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=960, bias=False)
        (bn2): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(960, 40, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(40, 960, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(960, 160, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (5): InvertedResidual(
        (conv_pw): Conv2d(160, 960, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(960, 960, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=960, bias=False)
        (bn2): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
```

```
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(960, 40, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(40, 960, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(960, 160, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (6): InvertedResidual(
        (conv_pw): Conv2d(160, 960, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(960, 960, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=960, bias=False)
        (bn2): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(960, 40, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(40, 960, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(960, 160, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
```

```
        )
        (drop_path): Identity()
      )
      (7): InvertedResidual(
        (conv_pw): Conv2d(160, 960, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(960, 960, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=960, bias=False)
        (bn2): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(960, 40, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(40, 960, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(960, 160, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (8): InvertedResidual(
        (conv_pw): Conv2d(160, 960, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(960, 960, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=960, bias=False)
        (bn2): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
```

```
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(960, 40, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(40, 960, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(960, 160, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (9): InvertedResidual(
        (conv_pw): Conv2d(160, 960, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(960, 960, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=960, bias=False)
        (bn2): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(960, 40, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(40, 960, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(960, 160, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
```

```
        (drop_path): Identity()
      )
    )
    (4): Sequential(
      (0): InvertedResidual(
        (conv_pw): Conv2d(160, 960, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(960, 960, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=960, bias=False)
        (bn2): BatchNormAct2d(
          960, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(960, 40, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(40, 960, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(960, 224, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          224, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (1): InvertedResidual(
        (conv_pw): Conv2d(224, 1344, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(1344, 1344, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1344, bias=False)
        (bn2): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
```

```
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(1344, 56, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(56, 1344, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(1344, 224, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          224, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (2): InvertedResidual(
        (conv_pw): Conv2d(224, 1344, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(1344, 1344, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1344, bias=False)
        (bn2): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(1344, 56, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(56, 1344, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(1344, 224, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          224, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
```

```
        )
        (drop_path): Identity()
      )
      (3): InvertedResidual(
        (conv_pw): Conv2d(224, 1344, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(1344, 1344, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1344, bias=False)
        (bn2): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(1344, 56, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(56, 1344, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(1344, 224, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          224, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (4): InvertedResidual(
        (conv_pw): Conv2d(224, 1344, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(1344, 1344, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1344, bias=False)
        (bn2): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
```

```
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(1344, 56, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(56, 1344, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(1344, 224, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          224, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (5): InvertedResidual(
        (conv_pw): Conv2d(224, 1344, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(1344, 1344, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1344, bias=False)
        (bn2): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(1344, 56, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(56, 1344, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(1344, 224, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          224, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
```

```
        (drop_path): Identity()
      )
      (6): InvertedResidual(
        (conv_pw): Conv2d(224, 1344, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(1344, 1344, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1344, bias=False)
        (bn2): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(1344, 56, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(56, 1344, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(1344, 224, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          224, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (7): InvertedResidual(
        (conv_pw): Conv2d(224, 1344, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(1344, 1344, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1344, bias=False)
        (bn2): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
```

```
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(1344, 56, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(56, 1344, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(1344, 224, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          224, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (8): InvertedResidual(
        (conv_pw): Conv2d(224, 1344, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(1344, 1344, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1344, bias=False)
        (bn2): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(1344, 56, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(56, 1344, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(1344, 224, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          224, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
```

```
      )
      (9): InvertedResidual(
        (conv_pw): Conv2d(224, 1344, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(1344, 1344, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1344, bias=False)
        (bn2): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(1344, 56, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(56, 1344, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(1344, 224, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          224, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
    )
    (5): Sequential(
      (0): InvertedResidual(
        (conv_pw): Conv2d(224, 1344, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2dSame(1344, 1344, kernel_size=(5, 5), stride=(2, 2),
groups=1344, bias=False)
        (bn2): BatchNormAct2d(
          1344, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
```

```
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(1344, 56, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(56, 1344, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(1344, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (1): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(2304, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
```

```
        (drop_path): Identity()
      )
      (2): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(2304, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (3): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
```

```
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(2304, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (4): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(2304, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
```

```
      )
      (5): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(2304, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (6): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
```

```
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(2304, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (7): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(2304, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
```

```
      (8): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(2304, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (9): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
```

```
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(2304, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (10): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(2304, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (11): InvertedResidual(
```

```
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(2304, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (12): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
```

```
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(2304, 384, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
    )
    (6): Sequential(
      (0): InvertedResidual(
        (conv_pw): Conv2d(384, 2304, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(2304, 2304, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=2304, bias=False)
        (bn2): BatchNormAct2d(
          2304, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(2304, 96, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(96, 2304, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(2304, 640, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          640, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
```

```
      (1): InvertedResidual(
        (conv_pw): Conv2d(640, 3840, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          3840, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(3840, 3840, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=3840, bias=False)
        (bn2): BatchNormAct2d(
          3840, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(3840, 160, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(160, 3840, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(3840, 640, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          640, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (2): InvertedResidual(
        (conv_pw): Conv2d(640, 3840, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          3840, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(3840, 3840, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=3840, bias=False)
        (bn2): BatchNormAct2d(
          3840, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
```

```
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(3840, 160, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(160, 3840, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(3840, 640, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          640, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
      (3): InvertedResidual(
        (conv_pw): Conv2d(640, 3840, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNormAct2d(
          3840, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (conv_dw): Conv2d(3840, 3840, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=3840, bias=False)
        (bn2): BatchNormAct2d(
          3840, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): SiLU(inplace=True)
        )
        (aa): Identity()
        (se): SqueezeExcite(
          (conv_reduce): Conv2d(3840, 160, kernel_size=(1, 1), stride=(1, 1))
          (act1): SiLU(inplace=True)
          (conv_expand): Conv2d(160, 3840, kernel_size=(1, 1), stride=(1, 1))
          (gate): Sigmoid()
        )
        (conv_pwl): Conv2d(3840, 640, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNormAct2d(
          640, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
          (drop): Identity()
          (act): Identity()
        )
        (drop_path): Identity()
      )
    )
```

```
  )
  (conv_head): Conv2d(640, 2560, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn2): BatchNormAct2d(
    2560, eps=0.001, momentum=0.1, affine=True, track_running_stats=True
    (drop): Identity()
    (act): SiLU(inplace=True)
  )
  (global_pool): SelectAdaptivePool2d(pool_type=avg,
flatten=Flatten(start_dim=1, end_dim=-1))
  (classifier): Sequential(
    (0): Linear(in_features=2560, out_features=1024, bias=True)
    (1): ReLU()
    (2): Dropout(p=0.4, inplace=False)
    (3): Linear(in_features=1024, out_features=256, bias=True)
    (4): ReLU()
    (5): Linear(in_features=256, out_features=2, bias=True)
  )
)
```

[39]:
```
from torchsummary import summary

summary(model, input_size=(3, config.img_size, config.img_size))
```

```
----------------------------------------------------------------
        Layer (type)               Output Shape         Param #
================================================================
        Conv2dSame-1         [-1, 64, 112, 112]           1,728
          Identity-2         [-1, 64, 112, 112]               0
              SiLU-3         [-1, 64, 112, 112]               0
    BatchNormAct2d-4         [-1, 64, 112, 112]             128
            Conv2d-5         [-1, 64, 112, 112]             576
          Identity-6         [-1, 64, 112, 112]               0
              SiLU-7         [-1, 64, 112, 112]               0
    BatchNormAct2d-8         [-1, 64, 112, 112]             128
          Identity-9         [-1, 64, 112, 112]               0
           Conv2d-10             [-1, 16, 1, 1]           1,040
             SiLU-11             [-1, 16, 1, 1]               0
           Conv2d-12             [-1, 64, 1, 1]           1,088
          Sigmoid-13             [-1, 64, 1, 1]               0
    SqueezeExcite-14         [-1, 64, 112, 112]               0
           Conv2d-15         [-1, 32, 112, 112]           2,048
         Identity-16         [-1, 32, 112, 112]               0
         Identity-17         [-1, 32, 112, 112]               0
   BatchNormAct2d-18         [-1, 32, 112, 112]              64
DepthwiseSeparableConv-19          [-1, 32, 112, 112]               0
           Conv2d-20         [-1, 32, 112, 112]             288
         Identity-21         [-1, 32, 112, 112]               0
```

```
            SiLU-22         [-1, 32, 112, 112]               0
   BatchNormAct2d-23         [-1, 32, 112, 112]              64
         Identity-24         [-1, 32, 112, 112]               0
           Conv2d-25              [-1, 8, 1, 1]             264
             SiLU-26              [-1, 8, 1, 1]               0
           Conv2d-27             [-1, 32, 1, 1]             288
          Sigmoid-28             [-1, 32, 1, 1]               0
    SqueezeExcite-29         [-1, 32, 112, 112]               0
           Conv2d-30         [-1, 32, 112, 112]           1,024
         Identity-31         [-1, 32, 112, 112]               0
         Identity-32         [-1, 32, 112, 112]               0
   BatchNormAct2d-33         [-1, 32, 112, 112]              64
         Identity-34         [-1, 32, 112, 112]               0
DepthwiseSeparableConv-35         [-1, 32, 112, 112]               0
           Conv2d-36         [-1, 32, 112, 112]             288
         Identity-37         [-1, 32, 112, 112]               0
             SiLU-38         [-1, 32, 112, 112]               0
   BatchNormAct2d-39         [-1, 32, 112, 112]              64
         Identity-40         [-1, 32, 112, 112]               0
           Conv2d-41              [-1, 8, 1, 1]             264
             SiLU-42              [-1, 8, 1, 1]               0
           Conv2d-43             [-1, 32, 1, 1]             288
          Sigmoid-44             [-1, 32, 1, 1]               0
    SqueezeExcite-45         [-1, 32, 112, 112]               0
           Conv2d-46         [-1, 32, 112, 112]           1,024
         Identity-47         [-1, 32, 112, 112]               0
         Identity-48         [-1, 32, 112, 112]               0
   BatchNormAct2d-49         [-1, 32, 112, 112]              64
         Identity-50         [-1, 32, 112, 112]               0
DepthwiseSeparableConv-51         [-1, 32, 112, 112]               0
           Conv2d-52         [-1, 32, 112, 112]             288
         Identity-53         [-1, 32, 112, 112]               0
             SiLU-54         [-1, 32, 112, 112]               0
   BatchNormAct2d-55         [-1, 32, 112, 112]              64
         Identity-56         [-1, 32, 112, 112]               0
           Conv2d-57              [-1, 8, 1, 1]             264
             SiLU-58              [-1, 8, 1, 1]               0
           Conv2d-59             [-1, 32, 1, 1]             288
          Sigmoid-60             [-1, 32, 1, 1]               0
    SqueezeExcite-61         [-1, 32, 112, 112]               0
           Conv2d-62         [-1, 32, 112, 112]           1,024
         Identity-63         [-1, 32, 112, 112]               0
         Identity-64         [-1, 32, 112, 112]               0
   BatchNormAct2d-65         [-1, 32, 112, 112]              64
         Identity-66         [-1, 32, 112, 112]               0
DepthwiseSeparableConv-67         [-1, 32, 112, 112]               0
           Conv2d-68        [-1, 192, 112, 112]           6,144
         Identity-69        [-1, 192, 112, 112]               0
```

SiLU-70	[-1, 192, 112, 112]	0
BatchNormAct2d-71	[-1, 192, 112, 112]	384
Conv2dSame-72	[-1, 192, 56, 56]	1,728
Identity-73	[-1, 192, 56, 56]	0
SiLU-74	[-1, 192, 56, 56]	0
BatchNormAct2d-75	[-1, 192, 56, 56]	384
Identity-76	[-1, 192, 56, 56]	0
Conv2d-77	[-1, 8, 1, 1]	1,544
SiLU-78	[-1, 8, 1, 1]	0
Conv2d-79	[-1, 192, 1, 1]	1,728
Sigmoid-80	[-1, 192, 1, 1]	0
SqueezeExcite-81	[-1, 192, 56, 56]	0
Conv2d-82	[-1, 48, 56, 56]	9,216
Identity-83	[-1, 48, 56, 56]	0
Identity-84	[-1, 48, 56, 56]	0
BatchNormAct2d-85	[-1, 48, 56, 56]	96
InvertedResidual-86	[-1, 48, 56, 56]	0
Conv2d-87	[-1, 288, 56, 56]	13,824
Identity-88	[-1, 288, 56, 56]	0
SiLU-89	[-1, 288, 56, 56]	0
BatchNormAct2d-90	[-1, 288, 56, 56]	576
Conv2d-91	[-1, 288, 56, 56]	2,592
Identity-92	[-1, 288, 56, 56]	0
SiLU-93	[-1, 288, 56, 56]	0
BatchNormAct2d-94	[-1, 288, 56, 56]	576
Identity-95	[-1, 288, 56, 56]	0
Conv2d-96	[-1, 12, 1, 1]	3,468
SiLU-97	[-1, 12, 1, 1]	0
Conv2d-98	[-1, 288, 1, 1]	3,744
Sigmoid-99	[-1, 288, 1, 1]	0
SqueezeExcite-100	[-1, 288, 56, 56]	0
Conv2d-101	[-1, 48, 56, 56]	13,824
Identity-102	[-1, 48, 56, 56]	0
Identity-103	[-1, 48, 56, 56]	0
BatchNormAct2d-104	[-1, 48, 56, 56]	96
Identity-105	[-1, 48, 56, 56]	0
InvertedResidual-106	[-1, 48, 56, 56]	0
Conv2d-107	[-1, 288, 56, 56]	13,824
Identity-108	[-1, 288, 56, 56]	0
SiLU-109	[-1, 288, 56, 56]	0
BatchNormAct2d-110	[-1, 288, 56, 56]	576
Conv2d-111	[-1, 288, 56, 56]	2,592
Identity-112	[-1, 288, 56, 56]	0
SiLU-113	[-1, 288, 56, 56]	0
BatchNormAct2d-114	[-1, 288, 56, 56]	576
Identity-115	[-1, 288, 56, 56]	0
Conv2d-116	[-1, 12, 1, 1]	3,468
SiLU-117	[-1, 12, 1, 1]	0

```
         Conv2d-118          [-1, 288, 1, 1]           3,744
        Sigmoid-119          [-1, 288, 1, 1]               0
  SqueezeExcite-120        [-1, 288, 56, 56]               0
         Conv2d-121         [-1, 48, 56, 56]          13,824
       Identity-122         [-1, 48, 56, 56]               0
       Identity-123         [-1, 48, 56, 56]               0
  BatchNormAct2d-124         [-1, 48, 56, 56]              96
       Identity-125         [-1, 48, 56, 56]               0
InvertedResidual-126         [-1, 48, 56, 56]               0
         Conv2d-127        [-1, 288, 56, 56]          13,824
       Identity-128        [-1, 288, 56, 56]               0
          SiLU-129        [-1, 288, 56, 56]               0
  BatchNormAct2d-130        [-1, 288, 56, 56]             576
         Conv2d-131        [-1, 288, 56, 56]           2,592
       Identity-132        [-1, 288, 56, 56]               0
          SiLU-133        [-1, 288, 56, 56]               0
  BatchNormAct2d-134        [-1, 288, 56, 56]             576
       Identity-135        [-1, 288, 56, 56]               0
         Conv2d-136           [-1, 12, 1, 1]           3,468
          SiLU-137           [-1, 12, 1, 1]               0
         Conv2d-138          [-1, 288, 1, 1]           3,744
        Sigmoid-139          [-1, 288, 1, 1]               0
  SqueezeExcite-140        [-1, 288, 56, 56]               0
         Conv2d-141         [-1, 48, 56, 56]          13,824
       Identity-142         [-1, 48, 56, 56]               0
       Identity-143         [-1, 48, 56, 56]               0
  BatchNormAct2d-144         [-1, 48, 56, 56]              96
       Identity-145         [-1, 48, 56, 56]               0
InvertedResidual-146         [-1, 48, 56, 56]               0
         Conv2d-147        [-1, 288, 56, 56]          13,824
       Identity-148        [-1, 288, 56, 56]               0
          SiLU-149        [-1, 288, 56, 56]               0
  BatchNormAct2d-150        [-1, 288, 56, 56]             576
         Conv2d-151        [-1, 288, 56, 56]           2,592
       Identity-152        [-1, 288, 56, 56]               0
          SiLU-153        [-1, 288, 56, 56]               0
  BatchNormAct2d-154        [-1, 288, 56, 56]             576
       Identity-155        [-1, 288, 56, 56]               0
         Conv2d-156           [-1, 12, 1, 1]           3,468
          SiLU-157           [-1, 12, 1, 1]               0
         Conv2d-158          [-1, 288, 1, 1]           3,744
        Sigmoid-159          [-1, 288, 1, 1]               0
  SqueezeExcite-160        [-1, 288, 56, 56]               0
         Conv2d-161         [-1, 48, 56, 56]          13,824
       Identity-162         [-1, 48, 56, 56]               0
       Identity-163         [-1, 48, 56, 56]               0
  BatchNormAct2d-164         [-1, 48, 56, 56]              96
       Identity-165         [-1, 48, 56, 56]               0
```

InvertedResidual-166	[-1, 48, 56, 56]	0
Conv2d-167	[-1, 288, 56, 56]	13,824
Identity-168	[-1, 288, 56, 56]	0
SiLU-169	[-1, 288, 56, 56]	0
BatchNormAct2d-170	[-1, 288, 56, 56]	576
Conv2d-171	[-1, 288, 56, 56]	2,592
Identity-172	[-1, 288, 56, 56]	0
SiLU-173	[-1, 288, 56, 56]	0
BatchNormAct2d-174	[-1, 288, 56, 56]	576
Identity-175	[-1, 288, 56, 56]	0
Conv2d-176	[-1, 12, 1, 1]	3,468
SiLU-177	[-1, 12, 1, 1]	0
Conv2d-178	[-1, 288, 1, 1]	3,744
Sigmoid-179	[-1, 288, 1, 1]	0
SqueezeExcite-180	[-1, 288, 56, 56]	0
Conv2d-181	[-1, 48, 56, 56]	13,824
Identity-182	[-1, 48, 56, 56]	0
Identity-183	[-1, 48, 56, 56]	0
BatchNormAct2d-184	[-1, 48, 56, 56]	96
Identity-185	[-1, 48, 56, 56]	0
InvertedResidual-186	[-1, 48, 56, 56]	0
Conv2d-187	[-1, 288, 56, 56]	13,824
Identity-188	[-1, 288, 56, 56]	0
SiLU-189	[-1, 288, 56, 56]	0
BatchNormAct2d-190	[-1, 288, 56, 56]	576
Conv2d-191	[-1, 288, 56, 56]	2,592
Identity-192	[-1, 288, 56, 56]	0
SiLU-193	[-1, 288, 56, 56]	0
BatchNormAct2d-194	[-1, 288, 56, 56]	576
Identity-195	[-1, 288, 56, 56]	0
Conv2d-196	[-1, 12, 1, 1]	3,468
SiLU-197	[-1, 12, 1, 1]	0
Conv2d-198	[-1, 288, 1, 1]	3,744
Sigmoid-199	[-1, 288, 1, 1]	0
SqueezeExcite-200	[-1, 288, 56, 56]	0
Conv2d-201	[-1, 48, 56, 56]	13,824
Identity-202	[-1, 48, 56, 56]	0
Identity-203	[-1, 48, 56, 56]	0
BatchNormAct2d-204	[-1, 48, 56, 56]	96
Identity-205	[-1, 48, 56, 56]	0
InvertedResidual-206	[-1, 48, 56, 56]	0
Conv2d-207	[-1, 288, 56, 56]	13,824
Identity-208	[-1, 288, 56, 56]	0
SiLU-209	[-1, 288, 56, 56]	0
BatchNormAct2d-210	[-1, 288, 56, 56]	576
Conv2dSame-211	[-1, 288, 28, 28]	7,200
Identity-212	[-1, 288, 28, 28]	0
SiLU-213	[-1, 288, 28, 28]	0

```
  BatchNormAct2d-214          [-1, 288, 28, 28]             576
        Identity-215          [-1, 288, 28, 28]               0
          Conv2d-216             [-1, 12, 1, 1]           3,468
            SiLU-217             [-1, 12, 1, 1]               0
          Conv2d-218            [-1, 288, 1, 1]           3,744
         Sigmoid-219            [-1, 288, 1, 1]               0
   SqueezeExcite-220          [-1, 288, 28, 28]               0
          Conv2d-221           [-1, 80, 28, 28]          23,040
        Identity-222           [-1, 80, 28, 28]               0
        Identity-223           [-1, 80, 28, 28]               0
  BatchNormAct2d-224           [-1, 80, 28, 28]             160
InvertedResidual-225           [-1, 80, 28, 28]               0
          Conv2d-226          [-1, 480, 28, 28]          38,400
        Identity-227          [-1, 480, 28, 28]               0
            SiLU-228          [-1, 480, 28, 28]               0
  BatchNormAct2d-229          [-1, 480, 28, 28]             960
          Conv2d-230          [-1, 480, 28, 28]          12,000
        Identity-231          [-1, 480, 28, 28]               0
            SiLU-232          [-1, 480, 28, 28]               0
  BatchNormAct2d-233          [-1, 480, 28, 28]             960
        Identity-234          [-1, 480, 28, 28]               0
          Conv2d-235             [-1, 20, 1, 1]           9,620
            SiLU-236             [-1, 20, 1, 1]               0
          Conv2d-237            [-1, 480, 1, 1]          10,080
         Sigmoid-238            [-1, 480, 1, 1]               0
   SqueezeExcite-239          [-1, 480, 28, 28]               0
          Conv2d-240           [-1, 80, 28, 28]          38,400
        Identity-241           [-1, 80, 28, 28]               0
        Identity-242           [-1, 80, 28, 28]               0
  BatchNormAct2d-243           [-1, 80, 28, 28]             160
        Identity-244           [-1, 80, 28, 28]               0
InvertedResidual-245           [-1, 80, 28, 28]               0
          Conv2d-246          [-1, 480, 28, 28]          38,400
        Identity-247          [-1, 480, 28, 28]               0
            SiLU-248          [-1, 480, 28, 28]               0
  BatchNormAct2d-249          [-1, 480, 28, 28]             960
          Conv2d-250          [-1, 480, 28, 28]          12,000
        Identity-251          [-1, 480, 28, 28]               0
            SiLU-252          [-1, 480, 28, 28]               0
  BatchNormAct2d-253          [-1, 480, 28, 28]             960
        Identity-254          [-1, 480, 28, 28]               0
          Conv2d-255             [-1, 20, 1, 1]           9,620
            SiLU-256             [-1, 20, 1, 1]               0
          Conv2d-257            [-1, 480, 1, 1]          10,080
         Sigmoid-258            [-1, 480, 1, 1]               0
   SqueezeExcite-259          [-1, 480, 28, 28]               0
          Conv2d-260           [-1, 80, 28, 28]          38,400
        Identity-261           [-1, 80, 28, 28]               0
```

```
       Identity-262          [-1, 80, 28, 28]               0
  BatchNormAct2d-263          [-1, 80, 28, 28]             160
       Identity-264          [-1, 80, 28, 28]               0
InvertedResidual-265          [-1, 80, 28, 28]               0
          Conv2d-266         [-1, 480, 28, 28]          38,400
       Identity-267         [-1, 480, 28, 28]               0
           SiLU-268         [-1, 480, 28, 28]               0
  BatchNormAct2d-269         [-1, 480, 28, 28]             960
          Conv2d-270         [-1, 480, 28, 28]          12,000
       Identity-271         [-1, 480, 28, 28]               0
           SiLU-272         [-1, 480, 28, 28]               0
  BatchNormAct2d-273         [-1, 480, 28, 28]             960
       Identity-274         [-1, 480, 28, 28]               0
          Conv2d-275            [-1, 20, 1, 1]           9,620
           SiLU-276            [-1, 20, 1, 1]               0
          Conv2d-277           [-1, 480, 1, 1]          10,080
         Sigmoid-278           [-1, 480, 1, 1]               0
   SqueezeExcite-279         [-1, 480, 28, 28]               0
          Conv2d-280          [-1, 80, 28, 28]          38,400
       Identity-281          [-1, 80, 28, 28]               0
       Identity-282          [-1, 80, 28, 28]               0
  BatchNormAct2d-283          [-1, 80, 28, 28]             160
       Identity-284          [-1, 80, 28, 28]               0
InvertedResidual-285          [-1, 80, 28, 28]               0
          Conv2d-286         [-1, 480, 28, 28]          38,400
       Identity-287         [-1, 480, 28, 28]               0
           SiLU-288         [-1, 480, 28, 28]               0
  BatchNormAct2d-289         [-1, 480, 28, 28]             960
          Conv2d-290         [-1, 480, 28, 28]          12,000
       Identity-291         [-1, 480, 28, 28]               0
           SiLU-292         [-1, 480, 28, 28]               0
  BatchNormAct2d-293         [-1, 480, 28, 28]             960
       Identity-294         [-1, 480, 28, 28]               0
          Conv2d-295            [-1, 20, 1, 1]           9,620
           SiLU-296            [-1, 20, 1, 1]               0
          Conv2d-297           [-1, 480, 1, 1]          10,080
         Sigmoid-298           [-1, 480, 1, 1]               0
   SqueezeExcite-299         [-1, 480, 28, 28]               0
          Conv2d-300          [-1, 80, 28, 28]          38,400
       Identity-301          [-1, 80, 28, 28]               0
       Identity-302          [-1, 80, 28, 28]               0
  BatchNormAct2d-303          [-1, 80, 28, 28]             160
       Identity-304          [-1, 80, 28, 28]               0
InvertedResidual-305          [-1, 80, 28, 28]               0
          Conv2d-306         [-1, 480, 28, 28]          38,400
       Identity-307         [-1, 480, 28, 28]               0
           SiLU-308         [-1, 480, 28, 28]               0
  BatchNormAct2d-309         [-1, 480, 28, 28]             960
```

```
         Conv2d-310          [-1, 480, 28, 28]          12,000
       Identity-311          [-1, 480, 28, 28]               0
          SiLU-312           [-1, 480, 28, 28]               0
  BatchNormAct2d-313          [-1, 480, 28, 28]             960
       Identity-314          [-1, 480, 28, 28]               0
         Conv2d-315             [-1, 20, 1, 1]           9,620
          SiLU-316             [-1, 20, 1, 1]               0
         Conv2d-317           [-1, 480, 1, 1]          10,080
       Sigmoid-318            [-1, 480, 1, 1]               0
   SqueezeExcite-319          [-1, 480, 28, 28]               0
         Conv2d-320           [-1, 80, 28, 28]          38,400
       Identity-321           [-1, 80, 28, 28]               0
       Identity-322           [-1, 80, 28, 28]               0
  BatchNormAct2d-323           [-1, 80, 28, 28]             160
       Identity-324           [-1, 80, 28, 28]               0
InvertedResidual-325           [-1, 80, 28, 28]               0
         Conv2d-326          [-1, 480, 28, 28]          38,400
       Identity-327          [-1, 480, 28, 28]               0
          SiLU-328           [-1, 480, 28, 28]               0
  BatchNormAct2d-329          [-1, 480, 28, 28]             960
         Conv2d-330          [-1, 480, 28, 28]          12,000
       Identity-331          [-1, 480, 28, 28]               0
          SiLU-332           [-1, 480, 28, 28]               0
  BatchNormAct2d-333          [-1, 480, 28, 28]             960
       Identity-334          [-1, 480, 28, 28]               0
         Conv2d-335             [-1, 20, 1, 1]           9,620
          SiLU-336             [-1, 20, 1, 1]               0
         Conv2d-337           [-1, 480, 1, 1]          10,080
       Sigmoid-338            [-1, 480, 1, 1]               0
   SqueezeExcite-339          [-1, 480, 28, 28]               0
         Conv2d-340           [-1, 80, 28, 28]          38,400
       Identity-341           [-1, 80, 28, 28]               0
       Identity-342           [-1, 80, 28, 28]               0
  BatchNormAct2d-343           [-1, 80, 28, 28]             160
       Identity-344           [-1, 80, 28, 28]               0
InvertedResidual-345           [-1, 80, 28, 28]               0
         Conv2d-346          [-1, 480, 28, 28]          38,400
       Identity-347          [-1, 480, 28, 28]               0
          SiLU-348           [-1, 480, 28, 28]               0
  BatchNormAct2d-349          [-1, 480, 28, 28]             960
     Conv2dSame-350          [-1, 480, 14, 14]           4,320
       Identity-351          [-1, 480, 14, 14]               0
          SiLU-352           [-1, 480, 14, 14]               0
  BatchNormAct2d-353          [-1, 480, 14, 14]             960
       Identity-354          [-1, 480, 14, 14]               0
         Conv2d-355             [-1, 20, 1, 1]           9,620
          SiLU-356             [-1, 20, 1, 1]               0
         Conv2d-357           [-1, 480, 1, 1]          10,080
```

Sigmoid-358	[-1, 480, 1, 1]	0
SqueezeExcite-359	[-1, 480, 14, 14]	0
Conv2d-360	[-1, 160, 14, 14]	76,800
Identity-361	[-1, 160, 14, 14]	0
Identity-362	[-1, 160, 14, 14]	0
BatchNormAct2d-363	[-1, 160, 14, 14]	320
InvertedResidual-364	[-1, 160, 14, 14]	0
Conv2d-365	[-1, 960, 14, 14]	153,600
Identity-366	[-1, 960, 14, 14]	0
SiLU-367	[-1, 960, 14, 14]	0
BatchNormAct2d-368	[-1, 960, 14, 14]	1,920
Conv2d-369	[-1, 960, 14, 14]	8,640
Identity-370	[-1, 960, 14, 14]	0
SiLU-371	[-1, 960, 14, 14]	0
BatchNormAct2d-372	[-1, 960, 14, 14]	1,920
Identity-373	[-1, 960, 14, 14]	0
Conv2d-374	[-1, 40, 1, 1]	38,440
SiLU-375	[-1, 40, 1, 1]	0
Conv2d-376	[-1, 960, 1, 1]	39,360
Sigmoid-377	[-1, 960, 1, 1]	0
SqueezeExcite-378	[-1, 960, 14, 14]	0
Conv2d-379	[-1, 160, 14, 14]	153,600
Identity-380	[-1, 160, 14, 14]	0
Identity-381	[-1, 160, 14, 14]	0
BatchNormAct2d-382	[-1, 160, 14, 14]	320
Identity-383	[-1, 160, 14, 14]	0
InvertedResidual-384	[-1, 160, 14, 14]	0
Conv2d-385	[-1, 960, 14, 14]	153,600
Identity-386	[-1, 960, 14, 14]	0
SiLU-387	[-1, 960, 14, 14]	0
BatchNormAct2d-388	[-1, 960, 14, 14]	1,920
Conv2d-389	[-1, 960, 14, 14]	8,640
Identity-390	[-1, 960, 14, 14]	0
SiLU-391	[-1, 960, 14, 14]	0
BatchNormAct2d-392	[-1, 960, 14, 14]	1,920
Identity-393	[-1, 960, 14, 14]	0
Conv2d-394	[-1, 40, 1, 1]	38,440
SiLU-395	[-1, 40, 1, 1]	0
Conv2d-396	[-1, 960, 1, 1]	39,360
Sigmoid-397	[-1, 960, 1, 1]	0
SqueezeExcite-398	[-1, 960, 14, 14]	0
Conv2d-399	[-1, 160, 14, 14]	153,600
Identity-400	[-1, 160, 14, 14]	0
Identity-401	[-1, 160, 14, 14]	0
BatchNormAct2d-402	[-1, 160, 14, 14]	320
Identity-403	[-1, 160, 14, 14]	0
InvertedResidual-404	[-1, 160, 14, 14]	0
Conv2d-405	[-1, 960, 14, 14]	153,600

```
      Identity-406          [-1, 960, 14, 14]               0
          SiLU-407          [-1, 960, 14, 14]               0
  BatchNormAct2d-408          [-1, 960, 14, 14]           1,920
        Conv2d-409          [-1, 960, 14, 14]           8,640
      Identity-410          [-1, 960, 14, 14]               0
          SiLU-411          [-1, 960, 14, 14]               0
  BatchNormAct2d-412          [-1, 960, 14, 14]           1,920
      Identity-413          [-1, 960, 14, 14]               0
        Conv2d-414             [-1, 40, 1, 1]          38,440
          SiLU-415             [-1, 40, 1, 1]               0
        Conv2d-416            [-1, 960, 1, 1]          39,360
       Sigmoid-417            [-1, 960, 1, 1]               0
   SqueezeExcite-418          [-1, 960, 14, 14]               0
        Conv2d-419          [-1, 160, 14, 14]         153,600
      Identity-420          [-1, 160, 14, 14]               0
      Identity-421          [-1, 160, 14, 14]               0
  BatchNormAct2d-422          [-1, 160, 14, 14]             320
      Identity-423          [-1, 160, 14, 14]               0
InvertedResidual-424          [-1, 160, 14, 14]               0
        Conv2d-425          [-1, 960, 14, 14]         153,600
      Identity-426          [-1, 960, 14, 14]               0
          SiLU-427          [-1, 960, 14, 14]               0
  BatchNormAct2d-428          [-1, 960, 14, 14]           1,920
        Conv2d-429          [-1, 960, 14, 14]           8,640
      Identity-430          [-1, 960, 14, 14]               0
          SiLU-431          [-1, 960, 14, 14]               0
  BatchNormAct2d-432          [-1, 960, 14, 14]           1,920
      Identity-433          [-1, 960, 14, 14]               0
        Conv2d-434             [-1, 40, 1, 1]          38,440
          SiLU-435             [-1, 40, 1, 1]               0
        Conv2d-436            [-1, 960, 1, 1]          39,360
       Sigmoid-437            [-1, 960, 1, 1]               0
   SqueezeExcite-438          [-1, 960, 14, 14]               0
        Conv2d-439          [-1, 160, 14, 14]         153,600
      Identity-440          [-1, 160, 14, 14]               0
      Identity-441          [-1, 160, 14, 14]               0
  BatchNormAct2d-442          [-1, 160, 14, 14]             320
      Identity-443          [-1, 160, 14, 14]               0
InvertedResidual-444          [-1, 160, 14, 14]               0
        Conv2d-445          [-1, 960, 14, 14]         153,600
      Identity-446          [-1, 960, 14, 14]               0
          SiLU-447          [-1, 960, 14, 14]               0
  BatchNormAct2d-448          [-1, 960, 14, 14]           1,920
        Conv2d-449          [-1, 960, 14, 14]           8,640
      Identity-450          [-1, 960, 14, 14]               0
          SiLU-451          [-1, 960, 14, 14]               0
  BatchNormAct2d-452          [-1, 960, 14, 14]           1,920
      Identity-453          [-1, 960, 14, 14]               0
```

Conv2d-454	[-1, 40, 1, 1]	38,440
SiLU-455	[-1, 40, 1, 1]	0
Conv2d-456	[-1, 960, 1, 1]	39,360
Sigmoid-457	[-1, 960, 1, 1]	0
SqueezeExcite-458	[-1, 960, 14, 14]	0
Conv2d-459	[-1, 160, 14, 14]	153,600
Identity-460	[-1, 160, 14, 14]	0
Identity-461	[-1, 160, 14, 14]	0
BatchNormAct2d-462	[-1, 160, 14, 14]	320
Identity-463	[-1, 160, 14, 14]	0
InvertedResidual-464	[-1, 160, 14, 14]	0
Conv2d-465	[-1, 960, 14, 14]	153,600
Identity-466	[-1, 960, 14, 14]	0
SiLU-467	[-1, 960, 14, 14]	0
BatchNormAct2d-468	[-1, 960, 14, 14]	1,920
Conv2d-469	[-1, 960, 14, 14]	8,640
Identity-470	[-1, 960, 14, 14]	0
SiLU-471	[-1, 960, 14, 14]	0
BatchNormAct2d-472	[-1, 960, 14, 14]	1,920
Identity-473	[-1, 960, 14, 14]	0
Conv2d-474	[-1, 40, 1, 1]	38,440
SiLU-475	[-1, 40, 1, 1]	0
Conv2d-476	[-1, 960, 1, 1]	39,360
Sigmoid-477	[-1, 960, 1, 1]	0
SqueezeExcite-478	[-1, 960, 14, 14]	0
Conv2d-479	[-1, 160, 14, 14]	153,600
Identity-480	[-1, 160, 14, 14]	0
Identity-481	[-1, 160, 14, 14]	0
BatchNormAct2d-482	[-1, 160, 14, 14]	320
Identity-483	[-1, 160, 14, 14]	0
InvertedResidual-484	[-1, 160, 14, 14]	0
Conv2d-485	[-1, 960, 14, 14]	153,600
Identity-486	[-1, 960, 14, 14]	0
SiLU-487	[-1, 960, 14, 14]	0
BatchNormAct2d-488	[-1, 960, 14, 14]	1,920
Conv2d-489	[-1, 960, 14, 14]	8,640
Identity-490	[-1, 960, 14, 14]	0
SiLU-491	[-1, 960, 14, 14]	0
BatchNormAct2d-492	[-1, 960, 14, 14]	1,920
Identity-493	[-1, 960, 14, 14]	0
Conv2d-494	[-1, 40, 1, 1]	38,440
SiLU-495	[-1, 40, 1, 1]	0
Conv2d-496	[-1, 960, 1, 1]	39,360
Sigmoid-497	[-1, 960, 1, 1]	0
SqueezeExcite-498	[-1, 960, 14, 14]	0
Conv2d-499	[-1, 160, 14, 14]	153,600
Identity-500	[-1, 160, 14, 14]	0
Identity-501	[-1, 160, 14, 14]	0

```
  BatchNormAct2d-502         [-1, 160, 14, 14]             320
        Identity-503         [-1, 160, 14, 14]               0
InvertedResidual-504         [-1, 160, 14, 14]               0
          Conv2d-505         [-1, 960, 14, 14]         153,600
        Identity-506         [-1, 960, 14, 14]               0
           SiLU-507         [-1, 960, 14, 14]               0
  BatchNormAct2d-508         [-1, 960, 14, 14]           1,920
          Conv2d-509         [-1, 960, 14, 14]           8,640
        Identity-510         [-1, 960, 14, 14]               0
           SiLU-511         [-1, 960, 14, 14]               0
  BatchNormAct2d-512         [-1, 960, 14, 14]           1,920
        Identity-513         [-1, 960, 14, 14]               0
          Conv2d-514            [-1, 40, 1, 1]          38,440
           SiLU-515            [-1, 40, 1, 1]               0
          Conv2d-516           [-1, 960, 1, 1]          39,360
         Sigmoid-517           [-1, 960, 1, 1]               0
   SqueezeExcite-518         [-1, 960, 14, 14]               0
          Conv2d-519         [-1, 160, 14, 14]         153,600
        Identity-520         [-1, 160, 14, 14]               0
        Identity-521         [-1, 160, 14, 14]               0
  BatchNormAct2d-522         [-1, 160, 14, 14]             320
        Identity-523         [-1, 160, 14, 14]               0
InvertedResidual-524         [-1, 160, 14, 14]               0
          Conv2d-525         [-1, 960, 14, 14]         153,600
        Identity-526         [-1, 960, 14, 14]               0
           SiLU-527         [-1, 960, 14, 14]               0
  BatchNormAct2d-528         [-1, 960, 14, 14]           1,920
          Conv2d-529         [-1, 960, 14, 14]           8,640
        Identity-530         [-1, 960, 14, 14]               0
           SiLU-531         [-1, 960, 14, 14]               0
  BatchNormAct2d-532         [-1, 960, 14, 14]           1,920
        Identity-533         [-1, 960, 14, 14]               0
          Conv2d-534            [-1, 40, 1, 1]          38,440
           SiLU-535            [-1, 40, 1, 1]               0
          Conv2d-536           [-1, 960, 1, 1]          39,360
         Sigmoid-537           [-1, 960, 1, 1]               0
   SqueezeExcite-538         [-1, 960, 14, 14]               0
          Conv2d-539         [-1, 160, 14, 14]         153,600
        Identity-540         [-1, 160, 14, 14]               0
        Identity-541         [-1, 160, 14, 14]               0
  BatchNormAct2d-542         [-1, 160, 14, 14]             320
        Identity-543         [-1, 160, 14, 14]               0
InvertedResidual-544         [-1, 160, 14, 14]               0
          Conv2d-545         [-1, 960, 14, 14]         153,600
        Identity-546         [-1, 960, 14, 14]               0
           SiLU-547         [-1, 960, 14, 14]               0
  BatchNormAct2d-548         [-1, 960, 14, 14]           1,920
          Conv2d-549         [-1, 960, 14, 14]          24,000
```

```
       Identity-550          [-1, 960, 14, 14]               0
          SiLU-551          [-1, 960, 14, 14]               0
  BatchNormAct2d-552          [-1, 960, 14, 14]           1,920
       Identity-553          [-1, 960, 14, 14]               0
         Conv2d-554             [-1, 40, 1, 1]          38,440
          SiLU-555             [-1, 40, 1, 1]               0
         Conv2d-556            [-1, 960, 1, 1]          39,360
        Sigmoid-557            [-1, 960, 1, 1]               0
   SqueezeExcite-558          [-1, 960, 14, 14]               0
         Conv2d-559          [-1, 224, 14, 14]         215,040
       Identity-560          [-1, 224, 14, 14]               0
       Identity-561          [-1, 224, 14, 14]               0
  BatchNormAct2d-562          [-1, 224, 14, 14]             448
InvertedResidual-563          [-1, 224, 14, 14]               0
         Conv2d-564         [-1, 1344, 14, 14]         301,056
       Identity-565         [-1, 1344, 14, 14]               0
          SiLU-566         [-1, 1344, 14, 14]               0
  BatchNormAct2d-567         [-1, 1344, 14, 14]           2,688
         Conv2d-568         [-1, 1344, 14, 14]          33,600
       Identity-569         [-1, 1344, 14, 14]               0
          SiLU-570         [-1, 1344, 14, 14]               0
  BatchNormAct2d-571         [-1, 1344, 14, 14]           2,688
       Identity-572         [-1, 1344, 14, 14]               0
         Conv2d-573             [-1, 56, 1, 1]          75,320
          SiLU-574             [-1, 56, 1, 1]               0
         Conv2d-575           [-1, 1344, 1, 1]          76,608
        Sigmoid-576           [-1, 1344, 1, 1]               0
   SqueezeExcite-577         [-1, 1344, 14, 14]               0
         Conv2d-578          [-1, 224, 14, 14]         301,056
       Identity-579          [-1, 224, 14, 14]               0
       Identity-580          [-1, 224, 14, 14]               0
  BatchNormAct2d-581          [-1, 224, 14, 14]             448
       Identity-582          [-1, 224, 14, 14]               0
InvertedResidual-583          [-1, 224, 14, 14]               0
         Conv2d-584         [-1, 1344, 14, 14]         301,056
       Identity-585         [-1, 1344, 14, 14]               0
          SiLU-586         [-1, 1344, 14, 14]               0
  BatchNormAct2d-587         [-1, 1344, 14, 14]           2,688
         Conv2d-588         [-1, 1344, 14, 14]          33,600
       Identity-589         [-1, 1344, 14, 14]               0
          SiLU-590         [-1, 1344, 14, 14]               0
  BatchNormAct2d-591         [-1, 1344, 14, 14]           2,688
       Identity-592         [-1, 1344, 14, 14]               0
         Conv2d-593             [-1, 56, 1, 1]          75,320
          SiLU-594             [-1, 56, 1, 1]               0
         Conv2d-595           [-1, 1344, 1, 1]          76,608
        Sigmoid-596           [-1, 1344, 1, 1]               0
   SqueezeExcite-597         [-1, 1344, 14, 14]               0
```

```
         Conv2d-598          [-1, 224, 14, 14]         301,056
       Identity-599          [-1, 224, 14, 14]               0
       Identity-600          [-1, 224, 14, 14]               0
  BatchNormAct2d-601          [-1, 224, 14, 14]             448
       Identity-602          [-1, 224, 14, 14]               0
InvertedResidual-603          [-1, 224, 14, 14]               0
         Conv2d-604         [-1, 1344, 14, 14]         301,056
       Identity-605         [-1, 1344, 14, 14]               0
          SiLU-606         [-1, 1344, 14, 14]               0
  BatchNormAct2d-607         [-1, 1344, 14, 14]           2,688
         Conv2d-608         [-1, 1344, 14, 14]          33,600
       Identity-609         [-1, 1344, 14, 14]               0
          SiLU-610         [-1, 1344, 14, 14]               0
  BatchNormAct2d-611         [-1, 1344, 14, 14]           2,688
       Identity-612         [-1, 1344, 14, 14]               0
         Conv2d-613             [-1, 56, 1, 1]          75,320
          SiLU-614             [-1, 56, 1, 1]               0
         Conv2d-615           [-1, 1344, 1, 1]          76,608
        Sigmoid-616           [-1, 1344, 1, 1]               0
   SqueezeExcite-617         [-1, 1344, 14, 14]               0
         Conv2d-618          [-1, 224, 14, 14]         301,056
       Identity-619          [-1, 224, 14, 14]               0
       Identity-620          [-1, 224, 14, 14]               0
  BatchNormAct2d-621          [-1, 224, 14, 14]             448
       Identity-622          [-1, 224, 14, 14]               0
InvertedResidual-623          [-1, 224, 14, 14]               0
         Conv2d-624         [-1, 1344, 14, 14]         301,056
       Identity-625         [-1, 1344, 14, 14]               0
          SiLU-626         [-1, 1344, 14, 14]               0
  BatchNormAct2d-627         [-1, 1344, 14, 14]           2,688
         Conv2d-628         [-1, 1344, 14, 14]          33,600
       Identity-629         [-1, 1344, 14, 14]               0
          SiLU-630         [-1, 1344, 14, 14]               0
  BatchNormAct2d-631         [-1, 1344, 14, 14]           2,688
       Identity-632         [-1, 1344, 14, 14]               0
         Conv2d-633             [-1, 56, 1, 1]          75,320
          SiLU-634             [-1, 56, 1, 1]               0
         Conv2d-635           [-1, 1344, 1, 1]          76,608
        Sigmoid-636           [-1, 1344, 1, 1]               0
   SqueezeExcite-637         [-1, 1344, 14, 14]               0
         Conv2d-638          [-1, 224, 14, 14]         301,056
       Identity-639          [-1, 224, 14, 14]               0
       Identity-640          [-1, 224, 14, 14]               0
  BatchNormAct2d-641          [-1, 224, 14, 14]             448
       Identity-642          [-1, 224, 14, 14]               0
InvertedResidual-643          [-1, 224, 14, 14]               0
         Conv2d-644         [-1, 1344, 14, 14]         301,056
       Identity-645         [-1, 1344, 14, 14]               0
```

```
           SiLU-646        [-1, 1344, 14, 14]               0
  BatchNormAct2d-647        [-1, 1344, 14, 14]           2,688
          Conv2d-648        [-1, 1344, 14, 14]          33,600
        Identity-649        [-1, 1344, 14, 14]               0
           SiLU-650        [-1, 1344, 14, 14]               0
  BatchNormAct2d-651        [-1, 1344, 14, 14]           2,688
        Identity-652        [-1, 1344, 14, 14]               0
          Conv2d-653            [-1, 56, 1, 1]          75,320
           SiLU-654            [-1, 56, 1, 1]               0
          Conv2d-655          [-1, 1344, 1, 1]          76,608
         Sigmoid-656          [-1, 1344, 1, 1]               0
   SqueezeExcite-657        [-1, 1344, 14, 14]               0
          Conv2d-658         [-1, 224, 14, 14]         301,056
        Identity-659         [-1, 224, 14, 14]               0
        Identity-660         [-1, 224, 14, 14]               0
  BatchNormAct2d-661         [-1, 224, 14, 14]             448
        Identity-662         [-1, 224, 14, 14]               0
InvertedResidual-663         [-1, 224, 14, 14]               0
          Conv2d-664        [-1, 1344, 14, 14]         301,056
        Identity-665        [-1, 1344, 14, 14]               0
           SiLU-666        [-1, 1344, 14, 14]               0
  BatchNormAct2d-667        [-1, 1344, 14, 14]           2,688
          Conv2d-668        [-1, 1344, 14, 14]          33,600
        Identity-669        [-1, 1344, 14, 14]               0
           SiLU-670        [-1, 1344, 14, 14]               0
  BatchNormAct2d-671        [-1, 1344, 14, 14]           2,688
        Identity-672        [-1, 1344, 14, 14]               0
          Conv2d-673            [-1, 56, 1, 1]          75,320
           SiLU-674            [-1, 56, 1, 1]               0
          Conv2d-675          [-1, 1344, 1, 1]          76,608
         Sigmoid-676          [-1, 1344, 1, 1]               0
   SqueezeExcite-677        [-1, 1344, 14, 14]               0
          Conv2d-678         [-1, 224, 14, 14]         301,056
        Identity-679         [-1, 224, 14, 14]               0
        Identity-680         [-1, 224, 14, 14]               0
  BatchNormAct2d-681         [-1, 224, 14, 14]             448
        Identity-682         [-1, 224, 14, 14]               0
InvertedResidual-683         [-1, 224, 14, 14]               0
          Conv2d-684        [-1, 1344, 14, 14]         301,056
        Identity-685        [-1, 1344, 14, 14]               0
           SiLU-686        [-1, 1344, 14, 14]               0
  BatchNormAct2d-687        [-1, 1344, 14, 14]           2,688
          Conv2d-688        [-1, 1344, 14, 14]          33,600
        Identity-689        [-1, 1344, 14, 14]               0
           SiLU-690        [-1, 1344, 14, 14]               0
  BatchNormAct2d-691        [-1, 1344, 14, 14]           2,688
        Identity-692        [-1, 1344, 14, 14]               0
          Conv2d-693            [-1, 56, 1, 1]          75,320
```

```
           SiLU-694            [-1, 56, 1, 1]               0
        Conv2d-695           [-1, 1344, 1, 1]          76,608
       Sigmoid-696           [-1, 1344, 1, 1]               0
  SqueezeExcite-697         [-1, 1344, 14, 14]               0
        Conv2d-698          [-1, 224, 14, 14]         301,056
      Identity-699          [-1, 224, 14, 14]               0
      Identity-700          [-1, 224, 14, 14]               0
 BatchNormAct2d-701          [-1, 224, 14, 14]             448
      Identity-702          [-1, 224, 14, 14]               0
InvertedResidual-703          [-1, 224, 14, 14]               0
        Conv2d-704         [-1, 1344, 14, 14]         301,056
      Identity-705         [-1, 1344, 14, 14]               0
           SiLU-706         [-1, 1344, 14, 14]               0
 BatchNormAct2d-707         [-1, 1344, 14, 14]           2,688
        Conv2d-708         [-1, 1344, 14, 14]          33,600
      Identity-709         [-1, 1344, 14, 14]               0
           SiLU-710         [-1, 1344, 14, 14]               0
 BatchNormAct2d-711         [-1, 1344, 14, 14]           2,688
      Identity-712         [-1, 1344, 14, 14]               0
        Conv2d-713            [-1, 56, 1, 1]          75,320
           SiLU-714            [-1, 56, 1, 1]               0
        Conv2d-715           [-1, 1344, 1, 1]          76,608
       Sigmoid-716           [-1, 1344, 1, 1]               0
  SqueezeExcite-717         [-1, 1344, 14, 14]               0
        Conv2d-718          [-1, 224, 14, 14]         301,056
      Identity-719          [-1, 224, 14, 14]               0
      Identity-720          [-1, 224, 14, 14]               0
 BatchNormAct2d-721          [-1, 224, 14, 14]             448
      Identity-722          [-1, 224, 14, 14]               0
InvertedResidual-723          [-1, 224, 14, 14]               0
        Conv2d-724         [-1, 1344, 14, 14]         301,056
      Identity-725         [-1, 1344, 14, 14]               0
           SiLU-726         [-1, 1344, 14, 14]               0
 BatchNormAct2d-727         [-1, 1344, 14, 14]           2,688
        Conv2d-728         [-1, 1344, 14, 14]          33,600
      Identity-729         [-1, 1344, 14, 14]               0
           SiLU-730         [-1, 1344, 14, 14]               0
 BatchNormAct2d-731         [-1, 1344, 14, 14]           2,688
      Identity-732         [-1, 1344, 14, 14]               0
        Conv2d-733            [-1, 56, 1, 1]          75,320
           SiLU-734            [-1, 56, 1, 1]               0
        Conv2d-735           [-1, 1344, 1, 1]          76,608
       Sigmoid-736           [-1, 1344, 1, 1]               0
  SqueezeExcite-737         [-1, 1344, 14, 14]               0
        Conv2d-738          [-1, 224, 14, 14]         301,056
      Identity-739          [-1, 224, 14, 14]               0
      Identity-740          [-1, 224, 14, 14]               0
 BatchNormAct2d-741          [-1, 224, 14, 14]             448
```

Identity-742	[-1, 224, 14, 14]	0
InvertedResidual-743	[-1, 224, 14, 14]	0
Conv2d-744	[-1, 1344, 14, 14]	301,056
Identity-745	[-1, 1344, 14, 14]	0
SiLU-746	[-1, 1344, 14, 14]	0
BatchNormAct2d-747	[-1, 1344, 14, 14]	2,688
Conv2dSame-748	[-1, 1344, 7, 7]	33,600
Identity-749	[-1, 1344, 7, 7]	0
SiLU-750	[-1, 1344, 7, 7]	0
BatchNormAct2d-751	[-1, 1344, 7, 7]	2,688
Identity-752	[-1, 1344, 7, 7]	0
Conv2d-753	[-1, 56, 1, 1]	75,320
SiLU-754	[-1, 56, 1, 1]	0
Conv2d-755	[-1, 1344, 1, 1]	76,608
Sigmoid-756	[-1, 1344, 1, 1]	0
SqueezeExcite-757	[-1, 1344, 7, 7]	0
Conv2d-758	[-1, 384, 7, 7]	516,096
Identity-759	[-1, 384, 7, 7]	0
Identity-760	[-1, 384, 7, 7]	0
BatchNormAct2d-761	[-1, 384, 7, 7]	768
InvertedResidual-762	[-1, 384, 7, 7]	0
Conv2d-763	[-1, 2304, 7, 7]	884,736
Identity-764	[-1, 2304, 7, 7]	0
SiLU-765	[-1, 2304, 7, 7]	0
BatchNormAct2d-766	[-1, 2304, 7, 7]	4,608
Conv2d-767	[-1, 2304, 7, 7]	57,600
Identity-768	[-1, 2304, 7, 7]	0
SiLU-769	[-1, 2304, 7, 7]	0
BatchNormAct2d-770	[-1, 2304, 7, 7]	4,608
Identity-771	[-1, 2304, 7, 7]	0
Conv2d-772	[-1, 96, 1, 1]	221,280
SiLU-773	[-1, 96, 1, 1]	0
Conv2d-774	[-1, 2304, 1, 1]	223,488
Sigmoid-775	[-1, 2304, 1, 1]	0
SqueezeExcite-776	[-1, 2304, 7, 7]	0
Conv2d-777	[-1, 384, 7, 7]	884,736
Identity-778	[-1, 384, 7, 7]	0
Identity-779	[-1, 384, 7, 7]	0
BatchNormAct2d-780	[-1, 384, 7, 7]	768
Identity-781	[-1, 384, 7, 7]	0
InvertedResidual-782	[-1, 384, 7, 7]	0
Conv2d-783	[-1, 2304, 7, 7]	884,736
Identity-784	[-1, 2304, 7, 7]	0
SiLU-785	[-1, 2304, 7, 7]	0
BatchNormAct2d-786	[-1, 2304, 7, 7]	4,608
Conv2d-787	[-1, 2304, 7, 7]	57,600
Identity-788	[-1, 2304, 7, 7]	0
SiLU-789	[-1, 2304, 7, 7]	0

```
  BatchNormAct2d-790          [-1, 2304, 7, 7]           4,608
        Identity-791          [-1, 2304, 7, 7]               0
          Conv2d-792            [-1, 96, 1, 1]         221,280
            SiLU-793            [-1, 96, 1, 1]               0
          Conv2d-794          [-1, 2304, 1, 1]         223,488
         Sigmoid-795          [-1, 2304, 1, 1]               0
   SqueezeExcite-796          [-1, 2304, 7, 7]               0
          Conv2d-797           [-1, 384, 7, 7]         884,736
        Identity-798           [-1, 384, 7, 7]               0
        Identity-799           [-1, 384, 7, 7]               0
  BatchNormAct2d-800           [-1, 384, 7, 7]             768
        Identity-801           [-1, 384, 7, 7]               0
InvertedResidual-802           [-1, 384, 7, 7]               0
          Conv2d-803          [-1, 2304, 7, 7]         884,736
        Identity-804          [-1, 2304, 7, 7]               0
            SiLU-805          [-1, 2304, 7, 7]               0
  BatchNormAct2d-806          [-1, 2304, 7, 7]           4,608
          Conv2d-807          [-1, 2304, 7, 7]          57,600
        Identity-808          [-1, 2304, 7, 7]               0
            SiLU-809          [-1, 2304, 7, 7]               0
  BatchNormAct2d-810          [-1, 2304, 7, 7]           4,608
        Identity-811          [-1, 2304, 7, 7]               0
          Conv2d-812            [-1, 96, 1, 1]         221,280
            SiLU-813            [-1, 96, 1, 1]               0
          Conv2d-814          [-1, 2304, 1, 1]         223,488
         Sigmoid-815          [-1, 2304, 1, 1]               0
   SqueezeExcite-816          [-1, 2304, 7, 7]               0
          Conv2d-817           [-1, 384, 7, 7]         884,736
        Identity-818           [-1, 384, 7, 7]               0
        Identity-819           [-1, 384, 7, 7]               0
  BatchNormAct2d-820           [-1, 384, 7, 7]             768
        Identity-821           [-1, 384, 7, 7]               0
InvertedResidual-822           [-1, 384, 7, 7]               0
          Conv2d-823          [-1, 2304, 7, 7]         884,736
        Identity-824          [-1, 2304, 7, 7]               0
            SiLU-825          [-1, 2304, 7, 7]               0
  BatchNormAct2d-826          [-1, 2304, 7, 7]           4,608
          Conv2d-827          [-1, 2304, 7, 7]          57,600
        Identity-828          [-1, 2304, 7, 7]               0
            SiLU-829          [-1, 2304, 7, 7]               0
  BatchNormAct2d-830          [-1, 2304, 7, 7]           4,608
        Identity-831          [-1, 2304, 7, 7]               0
          Conv2d-832            [-1, 96, 1, 1]         221,280
            SiLU-833            [-1, 96, 1, 1]               0
          Conv2d-834          [-1, 2304, 1, 1]         223,488
         Sigmoid-835          [-1, 2304, 1, 1]               0
   SqueezeExcite-836          [-1, 2304, 7, 7]               0
          Conv2d-837           [-1, 384, 7, 7]         884,736
```

Identity-838	[-1, 384, 7, 7]	0
Identity-839	[-1, 384, 7, 7]	0
BatchNormAct2d-840	[-1, 384, 7, 7]	768
Identity-841	[-1, 384, 7, 7]	0
InvertedResidual-842	[-1, 384, 7, 7]	0
Conv2d-843	[-1, 2304, 7, 7]	884,736
Identity-844	[-1, 2304, 7, 7]	0
SiLU-845	[-1, 2304, 7, 7]	0
BatchNormAct2d-846	[-1, 2304, 7, 7]	4,608
Conv2d-847	[-1, 2304, 7, 7]	57,600
Identity-848	[-1, 2304, 7, 7]	0
SiLU-849	[-1, 2304, 7, 7]	0
BatchNormAct2d-850	[-1, 2304, 7, 7]	4,608
Identity-851	[-1, 2304, 7, 7]	0
Conv2d-852	[-1, 96, 1, 1]	221,280
SiLU-853	[-1, 96, 1, 1]	0
Conv2d-854	[-1, 2304, 1, 1]	223,488
Sigmoid-855	[-1, 2304, 1, 1]	0
SqueezeExcite-856	[-1, 2304, 7, 7]	0
Conv2d-857	[-1, 384, 7, 7]	884,736
Identity-858	[-1, 384, 7, 7]	0
Identity-859	[-1, 384, 7, 7]	0
BatchNormAct2d-860	[-1, 384, 7, 7]	768
Identity-861	[-1, 384, 7, 7]	0
InvertedResidual-862	[-1, 384, 7, 7]	0
Conv2d-863	[-1, 2304, 7, 7]	884,736
Identity-864	[-1, 2304, 7, 7]	0
SiLU-865	[-1, 2304, 7, 7]	0
BatchNormAct2d-866	[-1, 2304, 7, 7]	4,608
Conv2d-867	[-1, 2304, 7, 7]	57,600
Identity-868	[-1, 2304, 7, 7]	0
SiLU-869	[-1, 2304, 7, 7]	0
BatchNormAct2d-870	[-1, 2304, 7, 7]	4,608
Identity-871	[-1, 2304, 7, 7]	0
Conv2d-872	[-1, 96, 1, 1]	221,280
SiLU-873	[-1, 96, 1, 1]	0
Conv2d-874	[-1, 2304, 1, 1]	223,488
Sigmoid-875	[-1, 2304, 1, 1]	0
SqueezeExcite-876	[-1, 2304, 7, 7]	0
Conv2d-877	[-1, 384, 7, 7]	884,736
Identity-878	[-1, 384, 7, 7]	0
Identity-879	[-1, 384, 7, 7]	0
BatchNormAct2d-880	[-1, 384, 7, 7]	768
Identity-881	[-1, 384, 7, 7]	0
InvertedResidual-882	[-1, 384, 7, 7]	0
Conv2d-883	[-1, 2304, 7, 7]	884,736
Identity-884	[-1, 2304, 7, 7]	0
SiLU-885	[-1, 2304, 7, 7]	0

```
  BatchNormAct2d-886          [-1, 2304, 7, 7]           4,608
          Conv2d-887          [-1, 2304, 7, 7]          57,600
       Identity-888           [-1, 2304, 7, 7]               0
          SiLU-889            [-1, 2304, 7, 7]               0
  BatchNormAct2d-890          [-1, 2304, 7, 7]           4,608
       Identity-891           [-1, 2304, 7, 7]               0
          Conv2d-892             [-1, 96, 1, 1]         221,280
          SiLU-893               [-1, 96, 1, 1]               0
          Conv2d-894           [-1, 2304, 1, 1]         223,488
         Sigmoid-895           [-1, 2304, 1, 1]               0
   SqueezeExcite-896           [-1, 2304, 7, 7]               0
          Conv2d-897            [-1, 384, 7, 7]         884,736
       Identity-898             [-1, 384, 7, 7]               0
       Identity-899             [-1, 384, 7, 7]               0
  BatchNormAct2d-900            [-1, 384, 7, 7]             768
       Identity-901             [-1, 384, 7, 7]               0
InvertedResidual-902            [-1, 384, 7, 7]               0
          Conv2d-903           [-1, 2304, 7, 7]         884,736
       Identity-904           [-1, 2304, 7, 7]               0
          SiLU-905            [-1, 2304, 7, 7]               0
  BatchNormAct2d-906          [-1, 2304, 7, 7]           4,608
          Conv2d-907          [-1, 2304, 7, 7]          57,600
       Identity-908           [-1, 2304, 7, 7]               0
          SiLU-909            [-1, 2304, 7, 7]               0
  BatchNormAct2d-910          [-1, 2304, 7, 7]           4,608
       Identity-911           [-1, 2304, 7, 7]               0
          Conv2d-912             [-1, 96, 1, 1]         221,280
          SiLU-913               [-1, 96, 1, 1]               0
          Conv2d-914           [-1, 2304, 1, 1]         223,488
         Sigmoid-915           [-1, 2304, 1, 1]               0
   SqueezeExcite-916           [-1, 2304, 7, 7]               0
          Conv2d-917            [-1, 384, 7, 7]         884,736
       Identity-918             [-1, 384, 7, 7]               0
       Identity-919             [-1, 384, 7, 7]               0
  BatchNormAct2d-920            [-1, 384, 7, 7]             768
       Identity-921             [-1, 384, 7, 7]               0
InvertedResidual-922            [-1, 384, 7, 7]               0
          Conv2d-923           [-1, 2304, 7, 7]         884,736
       Identity-924           [-1, 2304, 7, 7]               0
          SiLU-925            [-1, 2304, 7, 7]               0
  BatchNormAct2d-926          [-1, 2304, 7, 7]           4,608
          Conv2d-927          [-1, 2304, 7, 7]          57,600
       Identity-928           [-1, 2304, 7, 7]               0
          SiLU-929            [-1, 2304, 7, 7]               0
  BatchNormAct2d-930          [-1, 2304, 7, 7]           4,608
       Identity-931           [-1, 2304, 7, 7]               0
          Conv2d-932             [-1, 96, 1, 1]         221,280
          SiLU-933               [-1, 96, 1, 1]               0
```

Conv2d-934	[-1, 2304, 1, 1]	223,488
Sigmoid-935	[-1, 2304, 1, 1]	0
SqueezeExcite-936	[-1, 2304, 7, 7]	0
Conv2d-937	[-1, 384, 7, 7]	884,736
Identity-938	[-1, 384, 7, 7]	0
Identity-939	[-1, 384, 7, 7]	0
BatchNormAct2d-940	[-1, 384, 7, 7]	768
Identity-941	[-1, 384, 7, 7]	0
InvertedResidual-942	[-1, 384, 7, 7]	0
Conv2d-943	[-1, 2304, 7, 7]	884,736
Identity-944	[-1, 2304, 7, 7]	0
SiLU-945	[-1, 2304, 7, 7]	0
BatchNormAct2d-946	[-1, 2304, 7, 7]	4,608
Conv2d-947	[-1, 2304, 7, 7]	57,600
Identity-948	[-1, 2304, 7, 7]	0
SiLU-949	[-1, 2304, 7, 7]	0
BatchNormAct2d-950	[-1, 2304, 7, 7]	4,608
Identity-951	[-1, 2304, 7, 7]	0
Conv2d-952	[-1, 96, 1, 1]	221,280
SiLU-953	[-1, 96, 1, 1]	0
Conv2d-954	[-1, 2304, 1, 1]	223,488
Sigmoid-955	[-1, 2304, 1, 1]	0
SqueezeExcite-956	[-1, 2304, 7, 7]	0
Conv2d-957	[-1, 384, 7, 7]	884,736
Identity-958	[-1, 384, 7, 7]	0
Identity-959	[-1, 384, 7, 7]	0
BatchNormAct2d-960	[-1, 384, 7, 7]	768
Identity-961	[-1, 384, 7, 7]	0
InvertedResidual-962	[-1, 384, 7, 7]	0
Conv2d-963	[-1, 2304, 7, 7]	884,736
Identity-964	[-1, 2304, 7, 7]	0
SiLU-965	[-1, 2304, 7, 7]	0
BatchNormAct2d-966	[-1, 2304, 7, 7]	4,608
Conv2d-967	[-1, 2304, 7, 7]	57,600
Identity-968	[-1, 2304, 7, 7]	0
SiLU-969	[-1, 2304, 7, 7]	0
BatchNormAct2d-970	[-1, 2304, 7, 7]	4,608
Identity-971	[-1, 2304, 7, 7]	0
Conv2d-972	[-1, 96, 1, 1]	221,280
SiLU-973	[-1, 96, 1, 1]	0
Conv2d-974	[-1, 2304, 1, 1]	223,488
Sigmoid-975	[-1, 2304, 1, 1]	0
SqueezeExcite-976	[-1, 2304, 7, 7]	0
Conv2d-977	[-1, 384, 7, 7]	884,736
Identity-978	[-1, 384, 7, 7]	0
Identity-979	[-1, 384, 7, 7]	0
BatchNormAct2d-980	[-1, 384, 7, 7]	768
Identity-981	[-1, 384, 7, 7]	0

InvertedResidual-982	[-1, 384, 7, 7]	0
Conv2d-983	[-1, 2304, 7, 7]	884,736
Identity-984	[-1, 2304, 7, 7]	0
SiLU-985	[-1, 2304, 7, 7]	0
BatchNormAct2d-986	[-1, 2304, 7, 7]	4,608
Conv2d-987	[-1, 2304, 7, 7]	57,600
Identity-988	[-1, 2304, 7, 7]	0
SiLU-989	[-1, 2304, 7, 7]	0
BatchNormAct2d-990	[-1, 2304, 7, 7]	4,608
Identity-991	[-1, 2304, 7, 7]	0
Conv2d-992	[-1, 96, 1, 1]	221,280
SiLU-993	[-1, 96, 1, 1]	0
Conv2d-994	[-1, 2304, 1, 1]	223,488
Sigmoid-995	[-1, 2304, 1, 1]	0
SqueezeExcite-996	[-1, 2304, 7, 7]	0
Conv2d-997	[-1, 384, 7, 7]	884,736
Identity-998	[-1, 384, 7, 7]	0
Identity-999	[-1, 384, 7, 7]	0
BatchNormAct2d-1000	[-1, 384, 7, 7]	768
Identity-1001	[-1, 384, 7, 7]	0
InvertedResidual-1002	[-1, 384, 7, 7]	0
Conv2d-1003	[-1, 2304, 7, 7]	884,736
Identity-1004	[-1, 2304, 7, 7]	0
SiLU-1005	[-1, 2304, 7, 7]	0
BatchNormAct2d-1006	[-1, 2304, 7, 7]	4,608
Conv2d-1007	[-1, 2304, 7, 7]	20,736
Identity-1008	[-1, 2304, 7, 7]	0
SiLU-1009	[-1, 2304, 7, 7]	0
BatchNormAct2d-1010	[-1, 2304, 7, 7]	4,608
Identity-1011	[-1, 2304, 7, 7]	0
Conv2d-1012	[-1, 96, 1, 1]	221,280
SiLU-1013	[-1, 96, 1, 1]	0
Conv2d-1014	[-1, 2304, 1, 1]	223,488
Sigmoid-1015	[-1, 2304, 1, 1]	0
SqueezeExcite-1016	[-1, 2304, 7, 7]	0
Conv2d-1017	[-1, 640, 7, 7]	1,474,560
Identity-1018	[-1, 640, 7, 7]	0
Identity-1019	[-1, 640, 7, 7]	0
BatchNormAct2d-1020	[-1, 640, 7, 7]	1,280
InvertedResidual-1021	[-1, 640, 7, 7]	0
Conv2d-1022	[-1, 3840, 7, 7]	2,457,600
Identity-1023	[-1, 3840, 7, 7]	0
SiLU-1024	[-1, 3840, 7, 7]	0
BatchNormAct2d-1025	[-1, 3840, 7, 7]	7,680
Conv2d-1026	[-1, 3840, 7, 7]	34,560
Identity-1027	[-1, 3840, 7, 7]	0
SiLU-1028	[-1, 3840, 7, 7]	0
BatchNormAct2d-1029	[-1, 3840, 7, 7]	7,680

```
      Identity-1030          [-1, 3840, 7, 7]               0
        Conv2d-1031           [-1, 160, 1, 1]         614,560
          SiLU-1032           [-1, 160, 1, 1]               0
        Conv2d-1033          [-1, 3840, 1, 1]         618,240
       Sigmoid-1034          [-1, 3840, 1, 1]               0
  SqueezeExcite-1035          [-1, 3840, 7, 7]               0
        Conv2d-1036           [-1, 640, 7, 7]       2,457,600
      Identity-1037           [-1, 640, 7, 7]               0
      Identity-1038           [-1, 640, 7, 7]               0
 BatchNormAct2d-1039           [-1, 640, 7, 7]           1,280
      Identity-1040           [-1, 640, 7, 7]               0
InvertedResidual-1041           [-1, 640, 7, 7]               0
        Conv2d-1042          [-1, 3840, 7, 7]       2,457,600
      Identity-1043          [-1, 3840, 7, 7]               0
          SiLU-1044          [-1, 3840, 7, 7]               0
 BatchNormAct2d-1045          [-1, 3840, 7, 7]           7,680
        Conv2d-1046          [-1, 3840, 7, 7]          34,560
      Identity-1047          [-1, 3840, 7, 7]               0
          SiLU-1048          [-1, 3840, 7, 7]               0
 BatchNormAct2d-1049          [-1, 3840, 7, 7]           7,680
      Identity-1050          [-1, 3840, 7, 7]               0
        Conv2d-1051           [-1, 160, 1, 1]         614,560
          SiLU-1052           [-1, 160, 1, 1]               0
        Conv2d-1053          [-1, 3840, 1, 1]         618,240
       Sigmoid-1054          [-1, 3840, 1, 1]               0
  SqueezeExcite-1055          [-1, 3840, 7, 7]               0
        Conv2d-1056           [-1, 640, 7, 7]       2,457,600
      Identity-1057           [-1, 640, 7, 7]               0
      Identity-1058           [-1, 640, 7, 7]               0
 BatchNormAct2d-1059           [-1, 640, 7, 7]           1,280
      Identity-1060           [-1, 640, 7, 7]               0
InvertedResidual-1061           [-1, 640, 7, 7]               0
        Conv2d-1062          [-1, 3840, 7, 7]       2,457,600
      Identity-1063          [-1, 3840, 7, 7]               0
          SiLU-1064          [-1, 3840, 7, 7]               0
 BatchNormAct2d-1065          [-1, 3840, 7, 7]           7,680
        Conv2d-1066          [-1, 3840, 7, 7]          34,560
      Identity-1067          [-1, 3840, 7, 7]               0
          SiLU-1068          [-1, 3840, 7, 7]               0
 BatchNormAct2d-1069          [-1, 3840, 7, 7]           7,680
      Identity-1070          [-1, 3840, 7, 7]               0
        Conv2d-1071           [-1, 160, 1, 1]         614,560
          SiLU-1072           [-1, 160, 1, 1]               0
        Conv2d-1073          [-1, 3840, 1, 1]         618,240
       Sigmoid-1074          [-1, 3840, 1, 1]               0
  SqueezeExcite-1075          [-1, 3840, 7, 7]               0
        Conv2d-1076           [-1, 640, 7, 7]       2,457,600
      Identity-1077           [-1, 640, 7, 7]               0
```

```
       Identity-1078            [-1, 640, 7, 7]               0
 BatchNormAct2d-1079            [-1, 640, 7, 7]           1,280
       Identity-1080            [-1, 640, 7, 7]               0
InvertedResidual-1081            [-1, 640, 7, 7]               0
         Conv2d-1082           [-1, 2560, 7, 7]       1,638,400
       Identity-1083           [-1, 2560, 7, 7]               0
          SiLU-1084           [-1, 2560, 7, 7]               0
 BatchNormAct2d-1085           [-1, 2560, 7, 7]           5,120
AdaptiveAvgPool2d-1086             [-1, 2560, 1, 1]               0
        Flatten-1087                 [-1, 2560]               0
SelectAdaptivePool2d-1088                 [-1, 2560]               0
         Linear-1089                 [-1, 1024]       2,622,464
           ReLU-1090                 [-1, 1024]               0
        Dropout-1091                 [-1, 1024]               0
         Linear-1092                  [-1, 256]         262,400
           ReLU-1093                  [-1, 256]               0
         Linear-1094                    [-1, 2]             514
================================================================
Total params: 66,672,338
Trainable params: 2,885,378
Non-trainable params: 63,786,960
----------------------------------------------------------------
Input size (MB): 0.57
Forward/backward pass size (MB): 1539.97
Params size (MB): 254.33
Estimated Total Size (MB): 1794.87
----------------------------------------------------------------
```

1.14 Trainer class for training and validation

[40]:
```
# Define a named tuple for metrics
from collections import namedtuple

Metrics = namedtuple('Metrics', ['avg_loss', 'avg_accuracy', 'avg_recall',␣
 ↪'avg_precision', 'avg_f1_score'])
```

[41]:
```
class PneumoniaTrainer():

    def __init__(self, model, optimizer, criterion, device):
        self.model = model
        self.optimizer = optimizer
        self.criterion = criterion
        self.min_valid_loss = 1e10 # a large number
        self.best_valid_accuracy = 0.
        self.best_valid_recall = 0.
        self.best_valid_precision = 0.
        self.best_valid_f1_score = 0.
```

```
        self.device = device
        self.num_classes = 2

    def train_batch_loop(self, train_loader):

        self.model.to(self.device)
        self.model.train()

        total_loss = 0.

        metric_accuracy = torchmetrics.classification.
↪Accuracy(task="multiclass", num_classes=self.num_classes,␣
↪average="weighted").to(self.device)
        metric_recall = torchmetrics.classification.Recall(task="multiclass",␣
↪num_classes=self.num_classes, average="weighted").to(self.device)
        metric_precision = torchmetrics.classification.
↪Precision(task="multiclass", num_classes=self.num_classes,␣
↪average="weighted").to(self.device)
        metric_f1_score = torchmetrics.classification.
↪F1Score(task="multiclass", num_classes=self.num_classes, average="weighted").
↪to(self.device)

        for images, labels in tqdm(train_loader):
            images = images.to(self.device)
            labels = labels.to(self.device)

            logits = self.model(images)

            loss = self.criterion(logits, labels)

            self.optimizer.zero_grad()
            loss.backward()
            self.optimizer.step()

            total_loss += loss.item()

            accuracy_batch = metric_accuracy(logits, labels)
            recall_batch = metric_recall(logits, labels)
            precision_batch = metric_precision(logits, labels)
            f1_score_batch = metric_f1_score(logits, labels)


        avg_loss = total_loss / len(train_loader)
        avg_accuracy = metric_accuracy.compute().item()
        avg_recall = metric_recall.compute().item()
        avg_precision = metric_precision.compute().item()
        avg_f1_score = metric_f1_score.compute().item()
```

```
        metric_accuracy.reset()
        metric_recall.reset()
        metric_precision.reset()
        metric_f1_score.reset()

        return Metrics(avg_loss=avg_loss,
                       avg_accuracy=avg_accuracy,
                       avg_recall=avg_recall,
                       avg_precision=avg_precision,
                       avg_f1_score=avg_f1_score)

    def valid_batch_loop(self, valid_loader):

        self.model.to(self.device)
        self.model.eval() # turns off drop out and batch normalization layers

        total_loss = 0.

        metric_accuracy = torchmetrics.classification.
↪Accuracy(task="multiclass", num_classes=self.num_classes,␣
↪average="weighted").to(self.device)
        metric_recall = torchmetrics.classification.Recall(task="multiclass",␣
↪num_classes=self.num_classes, average="weighted").to(self.device)
        metric_precision = torchmetrics.classification.
↪Precision(task="multiclass", num_classes=self.num_classes,␣
↪average="weighted").to(self.device)
        metric_f1_score = torchmetrics.classification.
↪F1Score(task="multiclass", num_classes=self.num_classes, average="weighted").
↪to(self.device)

        with torch.no_grad():
            for images, labels in tqdm(valid_loader):
                images = images.to(self.device)
                labels = labels.to(self.device)

                logits = self.model(images)

                loss = self.criterion(logits, labels)

                total_loss += loss.item()

                accuracy_batch = metric_accuracy(logits, labels)
                recall_batch = metric_recall(logits, labels)
                precision_batch = metric_precision(logits, labels)
                f1_score_batch = metric_f1_score(logits, labels)
```

```
            avg_loss = total_loss / len(valid_loader)
            avg_accuracy = metric_accuracy.compute().item()
            avg_recall = metric_recall.compute().item()
            avg_precision = metric_precision.compute().item()
            avg_f1_score = metric_f1_score.compute().item()

            metric_accuracy.reset()
            metric_recall.reset()
            metric_precision.reset()
            metric_f1_score.reset()

        return Metrics(avg_loss=avg_loss,
                       avg_accuracy=avg_accuracy,
                       avg_recall=avg_recall,
                       avg_precision=avg_precision,
                       avg_f1_score=avg_f1_score)


    def fit(self, n_epoch, train_loader, valid_loader, save_filename):

        train_loss = []
        train_accuracy = []
        train_recall = []
        train_precision = []
        train_f1_score = []

        valid_loss = []
        valid_accuracy = []
        valid_recall = []
        valid_precision = []
        valid_f1_score = []

        # Early stopping parameters
        patience = 5
        patience_counter = 0

        for i_epoch in range(n_epoch):

            train_metrics = self.train_batch_loop(train_loader)

            valid_metrics = self.valid_batch_loop(valid_loader)

            train_loss.append(train_metrics.avg_loss)
            train_accuracy.append(train_metrics.avg_accuracy)
            train_recall.append(train_metrics.avg_recall)
            train_precision.append(train_metrics.avg_precision)
```

```
            train_f1_score.append(train_metrics.avg_f1_score)

            valid_loss.append(valid_metrics.avg_loss)
            valid_accuracy.append(valid_metrics.avg_accuracy)
            valid_recall.append(valid_metrics.avg_recall)
            valid_precision.append(valid_metrics.avg_precision)
            valid_f1_score.append(valid_metrics.avg_f1_score)

            print(f'Epoch: {i_epoch}, Train Loss: {train_metrics.avg_loss:.3f},␣
↪Train Acc.: {train_metrics.avg_accuracy:.3f}, ' +
                  f'Train recall: {train_metrics.avg_recall:.3f}, ' +
                  f'Train precision: {train_metrics.avg_precision:.3f}, Train␣
↪f1_score: {train_metrics.avg_f1_score:.3f}')

            print(f'Epoch: {i_epoch}, Valid Loss: {valid_metrics.avg_loss:.3f},␣
↪Valid Acc.: {valid_metrics.avg_accuracy:.3f}, ' +
                  f'Valid recall: {valid_metrics.avg_recall:.3f}, ' +
                  f'Valid precision: {valid_metrics.avg_precision:.3f}, Valid␣
↪f1_score: {valid_metrics.avg_f1_score:.3f}')

            if valid_metrics.avg_loss < self.min_valid_loss:
                self.min_valid_loss = valid_metrics.avg_loss
                self.best_valid_accuracy = valid_metrics.avg_accuracy
                self.best_valid_recall = valid_metrics.avg_recall
                self.best_valid_precision = valid_metrics.avg_precision
                self.best_valid_f1_score = valid_metrics.avg_f1_score
                patience_counter = 0
                torch.save(self.model.state_dict(), save_filename)
                print(f'Model saved to {save_filename}')
            else:
                patience_counter += 1
                if patience_counter >= patience:
                    print("Early stopping triggered.")
                    break

        print(f'\nBest Valid Loss: {self.min_valid_loss:.3f}')
        print(f'Best Valid Accuracy: {self.best_valid_accuracy:.3f}')
        print(f'Best Valid recall: {self.best_valid_recall:.3f}')
        print(f'Best Valid precision: {self.best_valid_precision:.3f}')
        print(f'Best Valid f1_score: {self.best_valid_f1_score:.3f}')

        train_lists = {'train_loss': train_loss,
                       'train_accuracy': train_accuracy,
                       'train_recall': train_recall,
                       'train_precision': train_precision,
                       'train_f1_score': train_f1_score
                       }
```

```
        valid_lists = {'valid_loss': valid_loss,
                       'valid_accuracy': valid_accuracy,
                       'valid_recall': valid_recall,
                       'valid_precision': valid_precision,
                       'valid_f1_score': valid_f1_score
                       }


        return train_lists, valid_lists
```

1.15 Train and fine-tune the models

[42]:
```
cnn_optimizer = torch.optim.Adam(cnn_model.parameters(), lr=config.lr_custom)
pretrained_optimizer = torch.optim.Adam(model.parameters(), lr=config.
 ↪lr_pretrained)
criterion = nn.CrossEntropyLoss(weight=class_weights.to(device))

cnn_trainer = PneumoniaTrainer(cnn_model, cnn_optimizer, criterion, device)
pretrained_trainer = PneumoniaTrainer(model, pretrained_optimizer, criterion,␣
 ↪device)

print('Training custom CNN model:')
n_epoch = 20
train_lists1, valid_lists1 = cnn_trainer.fit(n_epoch, train_loader_custom,␣
 ↪valid_loader_custom, 'pneumonia_cnn_model.pt')

print('Fine-tuning pretrained model:')
n_epoch = 20
save_filename_pretrained = 'pneumonia_pretrained_model_' + config.model_name +␣
 ↪'.pt'
train_lists2, valid_lists2  = pretrained_trainer.fit(n_epoch,␣
 ↪train_loader_pretrained, valid_loader_pretrained, save_filename_pretrained)
```

```
Training custom CNN model:
```

```
  0%|          | 0/41 [00:00<?, ?it/s]
```

```
  0%|          | 0/1 [00:00<?, ?it/s]
```

```
Epoch: 0, Train Loss: 0.573, Train Acc.: 0.677, Train recall: 0.677, Train
precision: 0.779, Train f1_score: 0.698
Epoch: 0, Valid Loss: 1.446, Valid Acc.: 0.500, Valid recall: 0.500, Valid
precision: 0.250, Valid f1_score: 0.333
Model saved to pneumonia_cnn_model.pt
```

```
  0%|          | 0/41 [00:00<?, ?it/s]
```

```
  0%|          | 0/1 [00:00<?, ?it/s]
```

```
Epoch: 1, Train Loss: 0.392, Train Acc.: 0.785, Train recall: 0.785, Train
precision: 0.867, Train f1_score: 0.798
Epoch: 1, Valid Loss: 0.588, Valid Acc.: 0.750, Valid recall: 0.750, Valid
precision: 0.767, Valid f1_score: 0.746
Model saved to pneumonia_cnn_model.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 2, Train Loss: 0.353, Train Acc.: 0.811, Train recall: 0.811, Train
precision: 0.879, Train f1_score: 0.823
Epoch: 2, Valid Loss: 0.558, Valid Acc.: 0.812, Valid recall: 0.812, Valid
precision: 0.864, Valid f1_score: 0.806
Model saved to pneumonia_cnn_model.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 3, Train Loss: 0.327, Train Acc.: 0.825, Train recall: 0.825, Train
precision: 0.886, Train f1_score: 0.835
Epoch: 3, Valid Loss: 0.530, Valid Acc.: 0.812, Valid recall: 0.812, Valid
precision: 0.864, Valid f1_score: 0.806
Model saved to pneumonia_cnn_model.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 4, Train Loss: 0.312, Train Acc.: 0.842, Train recall: 0.842, Train
precision: 0.896, Train f1_score: 0.851
Epoch: 4, Valid Loss: 0.514, Valid Acc.: 0.812, Valid recall: 0.812, Valid
precision: 0.864, Valid f1_score: 0.806
Model saved to pneumonia_cnn_model.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 5, Train Loss: 0.288, Train Acc.: 0.851, Train recall: 0.851, Train
precision: 0.899, Train f1_score: 0.859
Epoch: 5, Valid Loss: 0.456, Valid Acc.: 0.750, Valid recall: 0.750, Valid
precision: 0.767, Valid f1_score: 0.746
Model saved to pneumonia_cnn_model.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 6, Train Loss: 0.279, Train Acc.: 0.852, Train recall: 0.852, Train
precision: 0.899, Train f1_score: 0.860
Epoch: 6, Valid Loss: 0.430, Valid Acc.: 0.750, Valid recall: 0.750, Valid
precision: 0.767, Valid f1_score: 0.746
Model saved to pneumonia_cnn_model.pt
```

```
  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 7, Train Loss: 0.268, Train Acc.: 0.867, Train recall: 0.867, Train
precision: 0.907, Train f1_score: 0.874
Epoch: 7, Valid Loss: 0.421, Valid Acc.: 0.750, Valid recall: 0.750, Valid
precision: 0.767, Valid f1_score: 0.746
Model saved to pneumonia_cnn_model.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 8, Train Loss: 0.252, Train Acc.: 0.870, Train recall: 0.870, Train
precision: 0.908, Train f1_score: 0.877
Epoch: 8, Valid Loss: 0.401, Valid Acc.: 0.812, Valid recall: 0.812, Valid
precision: 0.817, Valid f1_score: 0.812
Model saved to pneumonia_cnn_model.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 9, Train Loss: 0.242, Train Acc.: 0.874, Train recall: 0.874, Train
precision: 0.911, Train f1_score: 0.880
Epoch: 9, Valid Loss: 0.405, Valid Acc.: 0.750, Valid recall: 0.750, Valid
precision: 0.767, Valid f1_score: 0.746

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 10, Train Loss: 0.235, Train Acc.: 0.880, Train recall: 0.880, Train
precision: 0.913, Train f1_score: 0.886
Epoch: 10, Valid Loss: 0.366, Valid Acc.: 0.812, Valid recall: 0.812, Valid
precision: 0.817, Valid f1_score: 0.812
Model saved to pneumonia_cnn_model.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 11, Train Loss: 0.228, Train Acc.: 0.887, Train recall: 0.887, Train
precision: 0.916, Train f1_score: 0.893
Epoch: 11, Valid Loss: 0.328, Valid Acc.: 0.938, Valid recall: 0.938, Valid
precision: 0.944, Valid f1_score: 0.937
Model saved to pneumonia_cnn_model.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 12, Train Loss: 0.217, Train Acc.: 0.889, Train recall: 0.889, Train
precision: 0.918, Train f1_score: 0.894
```

```
Epoch: 12, Valid Loss: 0.338, Valid Acc.: 0.875, Valid recall: 0.875, Valid
precision: 0.900, Valid f1_score: 0.873

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 13, Train Loss: 0.217, Train Acc.: 0.899, Train recall: 0.899, Train
precision: 0.923, Train f1_score: 0.904
Epoch: 13, Valid Loss: 0.324, Valid Acc.: 0.812, Valid recall: 0.812, Valid
precision: 0.817, Valid f1_score: 0.812
Model saved to pneumonia_cnn_model.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 14, Train Loss: 0.209, Train Acc.: 0.901, Train recall: 0.901, Train
precision: 0.924, Train f1_score: 0.905
Epoch: 14, Valid Loss: 0.310, Valid Acc.: 0.938, Valid recall: 0.938, Valid
precision: 0.944, Valid f1_score: 0.937
Model saved to pneumonia_cnn_model.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 15, Train Loss: 0.204, Train Acc.: 0.909, Train recall: 0.909, Train
precision: 0.929, Train f1_score: 0.913
Epoch: 15, Valid Loss: 0.319, Valid Acc.: 0.875, Valid recall: 0.875, Valid
precision: 0.875, Valid f1_score: 0.875

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 16, Train Loss: 0.192, Train Acc.: 0.915, Train recall: 0.915, Train
precision: 0.932, Train f1_score: 0.918
Epoch: 16, Valid Loss: 0.321, Valid Acc.: 0.875, Valid recall: 0.875, Valid
precision: 0.900, Valid f1_score: 0.873

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 17, Train Loss: 0.197, Train Acc.: 0.912, Train recall: 0.912, Train
precision: 0.930, Train f1_score: 0.915
Epoch: 17, Valid Loss: 0.346, Valid Acc.: 0.875, Valid recall: 0.875, Valid
precision: 0.900, Valid f1_score: 0.873

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 18, Train Loss: 0.185, Train Acc.: 0.917, Train recall: 0.917, Train
precision: 0.933, Train f1_score: 0.920
```

```
Epoch: 18, Valid Loss: 0.349, Valid Acc.: 0.875, Valid recall: 0.875, Valid
precision: 0.900, Valid f1_score: 0.873

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 19, Train Loss: 0.178, Train Acc.: 0.925, Train recall: 0.925, Train
precision: 0.938, Train f1_score: 0.928
Epoch: 19, Valid Loss: 0.337, Valid Acc.: 0.875, Valid recall: 0.875, Valid
precision: 0.900, Valid f1_score: 0.873
Early stopping triggered.

Best Valid Loss: 0.310
Best Valid Accuracy: 0.938
Best Valid recall: 0.938
Best Valid precision: 0.944
Best Valid f1_score: 0.937
Fine-tuning pretrained model:

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 0, Train Loss: 0.684, Train Acc.: 0.513, Train recall: 0.513, Train
precision: 0.798, Train f1_score: 0.519
Epoch: 0, Valid Loss: 0.663, Valid Acc.: 0.625, Valid recall: 0.625, Valid
precision: 0.786, Valid f1_score: 0.564
Model saved to pneumonia_pretrained_model_tf_efficientnet_b7.ap_in1k.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 1, Train Loss: 0.656, Train Acc.: 0.738, Train recall: 0.738, Train
precision: 0.866, Train f1_score: 0.755
Epoch: 1, Valid Loss: 0.618, Valid Acc.: 0.688, Valid recall: 0.688, Valid
precision: 0.808, Valid f1_score: 0.654
Model saved to pneumonia_pretrained_model_tf_efficientnet_b7.ap_in1k.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 2, Train Loss: 0.615, Train Acc.: 0.797, Train recall: 0.797, Train
precision: 0.882, Train f1_score: 0.809
Epoch: 2, Valid Loss: 0.557, Valid Acc.: 0.688, Valid recall: 0.688, Valid
precision: 0.808, Valid f1_score: 0.654
Model saved to pneumonia_pretrained_model_tf_efficientnet_b7.ap_in1k.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]
```

```
Epoch: 3, Train Loss: 0.553, Train Acc.: 0.833, Train recall: 0.833, Train
precision: 0.894, Train f1_score: 0.843
Epoch: 3, Valid Loss: 0.490, Valid Acc.: 0.688, Valid recall: 0.688, Valid
precision: 0.808, Valid f1_score: 0.654
Model saved to pneumonia_pretrained_model_tf_efficientnet_b7.ap_in1k.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 4, Train Loss: 0.481, Train Acc.: 0.850, Train recall: 0.850, Train
precision: 0.901, Train f1_score: 0.859
Epoch: 4, Valid Loss: 0.431, Valid Acc.: 0.750, Valid recall: 0.750, Valid
precision: 0.833, Valid f1_score: 0.733
Model saved to pneumonia_pretrained_model_tf_efficientnet_b7.ap_in1k.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 5, Train Loss: 0.410, Train Acc.: 0.865, Train recall: 0.865, Train
precision: 0.907, Train f1_score: 0.872
Epoch: 5, Valid Loss: 0.393, Valid Acc.: 0.750, Valid recall: 0.750, Valid
precision: 0.833, Valid f1_score: 0.733
Model saved to pneumonia_pretrained_model_tf_efficientnet_b7.ap_in1k.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 6, Train Loss: 0.348, Train Acc.: 0.876, Train recall: 0.876, Train
precision: 0.911, Train f1_score: 0.882
Epoch: 6, Valid Loss: 0.368, Valid Acc.: 0.750, Valid recall: 0.750, Valid
precision: 0.833, Valid f1_score: 0.733
Model saved to pneumonia_pretrained_model_tf_efficientnet_b7.ap_in1k.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 7, Train Loss: 0.303, Train Acc.: 0.891, Train recall: 0.891, Train
precision: 0.917, Train f1_score: 0.895
Epoch: 7, Valid Loss: 0.356, Valid Acc.: 0.750, Valid recall: 0.750, Valid
precision: 0.833, Valid f1_score: 0.733
Model saved to pneumonia_pretrained_model_tf_efficientnet_b7.ap_in1k.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 8, Train Loss: 0.275, Train Acc.: 0.898, Train recall: 0.898, Train
precision: 0.920, Train f1_score: 0.902
Epoch: 8, Valid Loss: 0.342, Valid Acc.: 0.750, Valid recall: 0.750, Valid
precision: 0.833, Valid f1_score: 0.733
Model saved to pneumonia_pretrained_model_tf_efficientnet_b7.ap_in1k.pt
```

```
  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 9, Train Loss: 0.249, Train Acc.: 0.900, Train recall: 0.900, Train
precision: 0.922, Train f1_score: 0.904
Epoch: 9, Valid Loss: 0.340, Valid Acc.: 0.750, Valid recall: 0.750, Valid
precision: 0.833, Valid f1_score: 0.733
Model saved to pneumonia_pretrained_model_tf_efficientnet_b7.ap_in1k.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 10, Train Loss: 0.228, Train Acc.: 0.906, Train recall: 0.906, Train
precision: 0.924, Train f1_score: 0.909
Epoch: 10, Valid Loss: 0.330, Valid Acc.: 0.812, Valid recall: 0.812, Valid
precision: 0.864, Valid f1_score: 0.806
Model saved to pneumonia_pretrained_model_tf_efficientnet_b7.ap_in1k.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 11, Train Loss: 0.212, Train Acc.: 0.913, Train recall: 0.913, Train
precision: 0.929, Train f1_score: 0.916
Epoch: 11, Valid Loss: 0.328, Valid Acc.: 0.812, Valid recall: 0.812, Valid
precision: 0.864, Valid f1_score: 0.806
Model saved to pneumonia_pretrained_model_tf_efficientnet_b7.ap_in1k.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 12, Train Loss: 0.200, Train Acc.: 0.922, Train recall: 0.922, Train
precision: 0.936, Train f1_score: 0.925
Epoch: 12, Valid Loss: 0.317, Valid Acc.: 0.812, Valid recall: 0.812, Valid
precision: 0.864, Valid f1_score: 0.806
Model saved to pneumonia_pretrained_model_tf_efficientnet_b7.ap_in1k.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 13, Train Loss: 0.196, Train Acc.: 0.921, Train recall: 0.921, Train
precision: 0.934, Train f1_score: 0.924
Epoch: 13, Valid Loss: 0.317, Valid Acc.: 0.812, Valid recall: 0.812, Valid
precision: 0.864, Valid f1_score: 0.806

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 14, Train Loss: 0.189, Train Acc.: 0.922, Train recall: 0.922, Train
precision: 0.934, Train f1_score: 0.924
Epoch: 14, Valid Loss: 0.315, Valid Acc.: 0.812, Valid recall: 0.812, Valid
```

```
precision: 0.864, Valid f1_score: 0.806
Model saved to pneumonia_pretrained_model_tf_efficientnet_b7.ap_in1k.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 15, Train Loss: 0.179, Train Acc.: 0.927, Train recall: 0.927, Train
precision: 0.938, Train f1_score: 0.929
Epoch: 15, Valid Loss: 0.308, Valid Acc.: 0.812, Valid recall: 0.812, Valid
precision: 0.864, Valid f1_score: 0.806
Model saved to pneumonia_pretrained_model_tf_efficientnet_b7.ap_in1k.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 16, Train Loss: 0.169, Train Acc.: 0.930, Train recall: 0.930, Train
precision: 0.940, Train f1_score: 0.932
Epoch: 16, Valid Loss: 0.310, Valid Acc.: 0.812, Valid recall: 0.812, Valid
precision: 0.864, Valid f1_score: 0.806

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 17, Train Loss: 0.161, Train Acc.: 0.937, Train recall: 0.937, Train
precision: 0.946, Train f1_score: 0.939
Epoch: 17, Valid Loss: 0.305, Valid Acc.: 0.812, Valid recall: 0.812, Valid
precision: 0.864, Valid f1_score: 0.806
Model saved to pneumonia_pretrained_model_tf_efficientnet_b7.ap_in1k.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 18, Train Loss: 0.161, Train Acc.: 0.936, Train recall: 0.936, Train
precision: 0.944, Train f1_score: 0.938
Epoch: 18, Valid Loss: 0.299, Valid Acc.: 0.812, Valid recall: 0.812, Valid
precision: 0.864, Valid f1_score: 0.806
Model saved to pneumonia_pretrained_model_tf_efficientnet_b7.ap_in1k.pt

  0%|          | 0/41 [00:00<?, ?it/s]

  0%|          | 0/1 [00:00<?, ?it/s]

Epoch: 19, Train Loss: 0.162, Train Acc.: 0.933, Train recall: 0.933, Train
precision: 0.941, Train f1_score: 0.934
Epoch: 19, Valid Loss: 0.309, Valid Acc.: 0.812, Valid recall: 0.812, Valid
precision: 0.864, Valid f1_score: 0.806


Best Valid Loss: 0.299
Best Valid Accuracy: 0.812
Best Valid recall: 0.812
```

```
Best Valid precision: 0.864
Best Valid f1_score: 0.806
```

1.16 Plot the histories

[43]:
```
fig, ax = plt.subplots(nrows=1, ncols=3, figsize=(15,4), dpi=150)
plt.suptitle("Custom CNN Model")
ax[0].plot(train_lists1['train_loss'], label='train_loss')
ax[0].plot(valid_lists1['valid_loss'], label='valid_loss')
ax[0].set_title('Loss History')
ax[0].set_xlabel('Epoch')
ax[0].set_ylabel('Loss')
ax[0].legend(loc='best')
ax[0].grid(True)
ax[0].set_xlim([-0.5, n_epoch-0.5])
ax[0].set_ylim([-0.02, 0.82])
ax[0].set_xticks(np.arange(0, n_epoch, step=2))

ax[1].plot(train_lists1['train_accuracy'], label='train_accuracy')
ax[1].plot(valid_lists1['valid_accuracy'], label='valid_accuracy')
ax[1].set_title('Weighted Accuracy History')
ax[1].set_xlabel('Epoch')
ax[1].set_ylabel('Weighted Accuracy')
ax[1].legend(loc='best')
ax[1].grid(True)
ax[1].set_xlim([-0.5, n_epoch-0.5])
ax[1].set_ylim([0.48, 1.02])
ax[1].set_xticks(np.arange(0, n_epoch, step=2))

ax[2].plot(train_lists1['train_recall'], label='train_recall')
ax[2].plot(valid_lists1['valid_recall'], label='valid_recall')
ax[2].set_title('Weighted Recall History')
ax[2].set_xlabel('Epoch')
ax[2].set_ylabel('Weighted Recall')
ax[2].legend(loc='best')
ax[2].grid(True)
ax[2].set_xlim([-0.5, n_epoch-0.5])
ax[2].set_ylim([0.48, 1.02])
ax[2].set_xticks(np.arange(0, n_epoch, step=2))

fig.tight_layout();

fig_filename = 'pneumonia_custom_model_' + f'_lr_{config.lr_custom}_bs_{config.
 ↪batch_size}' + '.png'
fig.savefig(fig_filename)
```

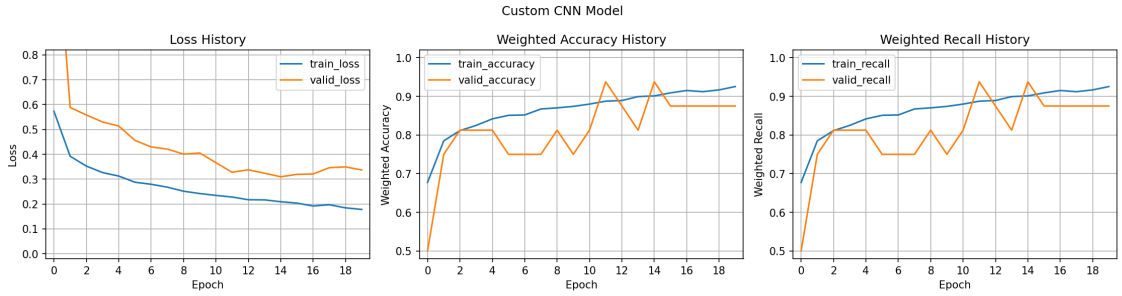

[44]:
```
fig, ax = plt.subplots(nrows=1, ncols=3, figsize=(15,4), dpi=150)
plt.suptitle("Fine-tuned Model")
ax[0].plot(train_lists2['train_loss'], label='train_loss')
ax[0].plot(valid_lists2['valid_loss'], label='valid_loss')
ax[0].set_title('Loss History')
ax[0].set_xlabel('Epoch')
ax[0].set_ylabel('Loss')
ax[0].legend(loc='best')
ax[0].grid(True)
ax[0].set_xlim([-0.5, n_epoch-0.5])
ax[0].set_ylim([-0.02, 0.82])
ax[0].set_xticks(np.arange(0, n_epoch, step=2))

ax[1].plot(train_lists2['train_accuracy'], label='train_accuracy')
ax[1].plot(valid_lists2['valid_accuracy'], label='valid_accuracy')
ax[1].set_title('Weighted Accuracy History')
ax[1].set_xlabel('Epoch')
ax[1].set_ylabel('Weighted Accuracy')
ax[1].legend(loc='best')
ax[1].grid(True)
ax[1].set_xlim([-0.5, n_epoch-0.5])
ax[1].set_ylim([0.48, 1.02])
ax[1].set_xticks(np.arange(0, n_epoch, step=2))

ax[2].plot(train_lists2['train_recall'], label='train_recall')
ax[2].plot(valid_lists2['valid_recall'], label='valid_recall')
ax[2].set_title('Weighted Recall History')
ax[2].set_xlabel('Epoch')
ax[2].set_ylabel('Weighted Recall')
ax[2].legend(loc='best')
ax[2].grid(True)
ax[2].set_xlim([-0.5, n_epoch-0.5])
ax[2].set_ylim([0.48, 1.02])
ax[2].set_xticks(np.arange(0, n_epoch, step=2))

fig.tight_layout();
```

```
fig_filename = 'pneumonia_pretrained_model_' + config.model_name +␣
  ↪f'_lr_{config.lr_pretrained}_bs_{config.batch_size}' + '.png'
fig.savefig(fig_filename)
```

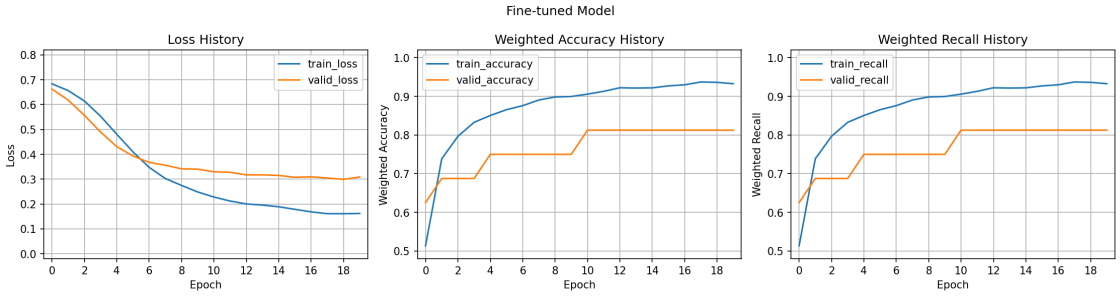

1.17 Evaluate the performance on testset

[45]:
```
status = cnn_model.load_state_dict(torch.load('pneumonia_cnn_model.pt',␣
  ↪map_location=device, weights_only=True))
print(f"Status: {status}")

model_size = os.path.getsize('pneumonia_cnn_model.pt') / (1024 * 1024)  # in MB
print(f"Custom CNN model size: {model_size:.2f} MB")

test_metrics = cnn_trainer.valid_batch_loop(test_loader_custom)

print(f"Test Loss: {test_metrics.avg_loss:.3f}")
print(f"Test Acc.: {test_metrics.avg_accuracy:.3f}")
print(f"Test Recall: {test_metrics.avg_recall:.3f}")
print(f"Test Precision: {test_metrics.avg_precision:.3f}")
print(f"Test F1-score: {test_metrics.avg_f1_score:.3f}")
```

```
Status: <All keys matched successfully>
Custom CNN model size: 12.66 MB
```

```
  0%|          | 0/5 [00:00<?, ?it/s]
```

```
Test Loss: 0.317
Test Acc.: 0.894
Test Recall: 0.894
Test Precision: 0.896
Test F1-score: 0.895
```

[46]:
```
save_filename_pretrained = 'pneumonia_pretrained_model_' + config.model_name +␣
  ↪'.pt'
```

```
status = model.load_state_dict(torch.load('/content/' +␣
 ↪save_filename_pretrained, map_location=device, weights_only=True))
print(f"Status: {status}")
print(f"save_filename_pretrained: {save_filename_pretrained}")

model_size = os.path.getsize(save_filename_pretrained) / (1024 * 1024)  # in MB
print(f"Pretrained model size: {model_size:.2f} MB")

test_metrics_pretrained = pretrained_trainer.
 ↪valid_batch_loop(test_loader_pretrained)

print(f"Test Loss: {test_metrics_pretrained.avg_loss:.3f}")
print(f"Test Acc.: {test_metrics_pretrained.avg_accuracy:.3f}")
print(f"Test Recall: {test_metrics_pretrained.avg_recall:.3f}")
print(f"Test Precision: {test_metrics_pretrained.avg_precision:.3f}")
print(f"Test F1-score: {test_metrics_pretrained.avg_f1_score:.3f}")
```

```
Status: <All keys matched successfully>
save_filename_pretrained:
pneumonia_pretrained_model_tf_efficientnet_b7.ap_in1k.pt
Pretrained model size: 255.98 MB
```

```
  0%|          | 0/5 [00:00<?, ?it/s]
```

```
Test Loss: 0.336
Test Acc.: 0.872
Test Recall: 0.872
Test Precision: 0.876
Test F1-score: 0.873
```

1.18 Show the prediction for a sample image

[47]:
```
def show_predicted_class(image, label_prob, label, mean, std, title):

    class_names = ['Normal', 'Pneumonia']

    label_prob = label_prob.cpu().detach().numpy().squeeze(0)

    image = image.to('cpu')
    image = image.permute(1, 2, 0)
    mean = torch.FloatTensor(mean)
    std = torch.FloatTensor(std)
    image = image*std + mean
    img = torch.clip(image, 0, 1)
    img = img.numpy()

    fig, (ax1, ax2) = plt.subplots(figsize=(6,3), ncols=2, dpi=150)
    plt.suptitle(title)
```

```
    ax1.imshow(img, cmap="gray")
    ax1.set_title(f'Ground Truth: {class_names[label]}')
    ax1.axis(False)

    ax2.barh(class_names, label_prob)
    ax2.set_aspect(0.2)
    ax2.set_yticks(class_names)
    ax2.set_yticklabels(class_names)
    ax2.set_title(f'Predicted Class: {class_names[label_prob.argmax()]}')
    ax2.set_xlim(0, 1.05)

    plt.tight_layout()
```

[48]:
```
import random

random_int = random.randint(0, len(testset_custom)-1)
# For custom model
image, label = testset_custom[random_int]
image = image.to(device)

logit = cnn_model(image.unsqueeze(0))
label_predict = F.softmax(logit, dim = 1)

show_predicted_class(image, label_predict, label, config.pn_mean, config.
 ↪pn_std, title="Custom CNN Model")

# For fine-tuned model
image, label = testset_pretrained[random_int]
image = image.to(device)

logit = model(image.unsqueeze(0))
label_predict = F.softmax(logit, dim = 1)

show_predicted_class(image, label_predict, label, config.in_mean, config.
 ↪in_std, title="Pretrained Model")
```

Custom CNN Model

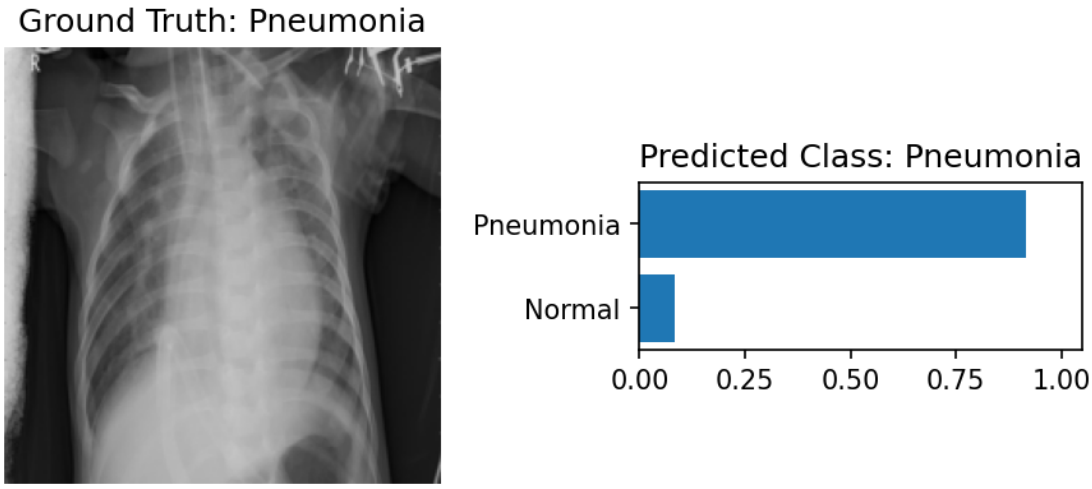

Pretrained Model

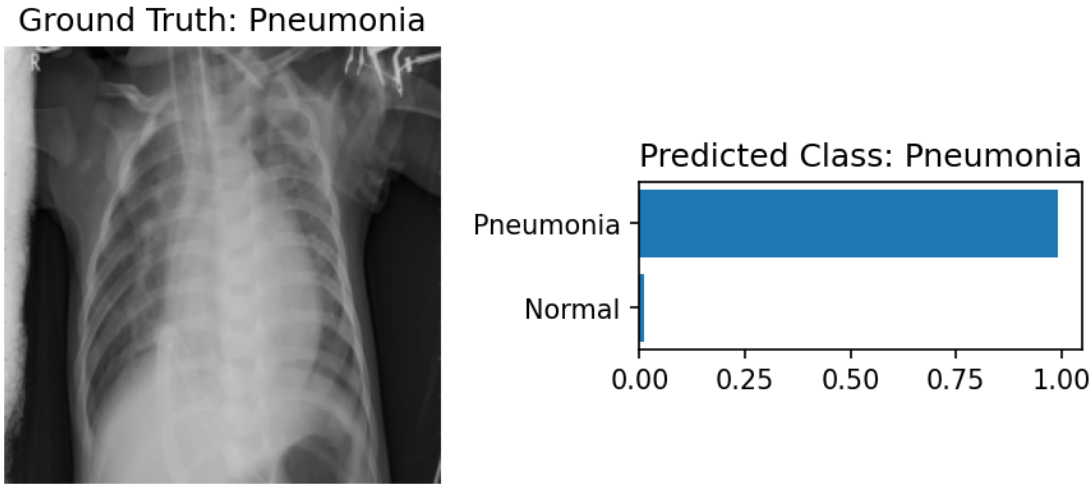

[48]:

[48]: